Fine-tuning collapses differentiation gradients in single-cell foundation models in a domain-specific manner

Kunhui Xu[a], Qi Dai[a,*]

[a]*National Clinical Research Center for Ocular Diseases, Eye Hospital, Wenzhou Medical University, Wenzhou, Zhejiang 325035, PR China*

Abstract

Supervised fine-tuning of single-cell foundation models is routine practice, yet whether it preserves the continuous differentiation geometry in pretrained representations has not been quantified. Using the human corneal epithelium, a four-stage differentiation continuum (limbal stem cells, transit-amplifying cells, basal, and differentiated) with known ordinal structure, we show that classification fine-tuning of Geneformer-V2-104M disrupts the differentiation gradient domain-specifically: catastrophic collapse on same-corpus data (Spearman ρ 0.806 $\rightarrow -0.008$) with preserved layer topology, mild collapse on truly out-of-corpus data (0.707 $\rightarrow$ 0.641), and improvement on cross-tissue data (0.24 $\rightarrow$ 0.46). An ordinal head (CORN) is the most geometry-faithful recipe ($\rho = 0.833$ at matched accuracy), preserved only within-corpus; of five alternative objectives, three failed outright, one preserved only partially (three-seed $\rho = 0.208$), and one was unstable across seeds. Primary claims survive BH-FDR control (26 comparisons) and label/feature-shuffle nulls. We release a contamination-aware benchmark pairing same-corpus and out-of-corpus external tests.

Keywords: single-cell foundation models, classification fine-tuning, differentiation gradient, ordinal regression, contamination-aware benchmark, corneal epithelium

1. Introduction

The corneal epithelium is maintained by a tightly regulated differentiation continuum from quiescent limbal stem cells (LSCs) at the limbal palisades of Vogt, through transit-amplifying cells (TACs), to committed basal epithelial cells, and ultimately to terminally differentiated squamous cells forming the protective ocular surface barrier [1]. This ordered progression is functionally required: disruption at any stage (LSC deficiency, aberrant amplification, or premature terminal differentiation) compromises corneal transparency and may lead to irreversible vision loss [2].

Single-cell RNA sequencing (scRNA-seq) has resolved this continuum at single-cell resolution, with marker genes for LSCs (e.g., *TCF4*, *POSTN*, *ABCG2*, and *TSPAN7* [3]; *TP63* and *KRT15* [4, 5]), TACs (e.g., *UBE2C*, *CDC20*, and *CENPF* [6]; *MKI67* and *BIRC5* [3]), and differentiated populations (e.g., *KRT12* and *KRT3* [4, 7, 8]) identified across independent

*Corresponding author

Email address: dq@mail.eye.ac.cn (Qi Dai)

studies and aggregated in community atlases such as the human ocular surface atlas in CELLxGENE Discover [9]. Yet these efforts produce study-specific embeddings that are not directly comparable across datasets, and no unified reference model maps individual cells to a shared, ordered differentiation-state space. The closest approach, a corneal cell-state reference model with support-vector-machine classification (cPredictor) [10], provides a supervised classifier rather than an evaluation of foundation-model representations or their geometric structure.

Foundation models pretrained on massive single-cell corpora offer a potential solution. Geneformer, a 12-layer Transformer pretrained on approximately 104 million cells from 2,903 datasets [11, 12], and related architectures such as scGPT [13] and scBERT [14], learn contextual gene–gene representations that transfer across tissues. Systematic evaluations have tempered expectations: a unified benchmark of thirteen single-cell foundation models found classical PCA remains competitive in several settings, with pretrained embeddings advantageous mainly in low-label and transfer-learning scenarios [15]; biology-driven reanalyses questioned the pre-train-then-fine-tune paradigm [16]; and large-scale evaluations found that, for cell-type annotation, frozen or zero-shot usage generally underperforms fine-tuning, and freezing pretrained layers was not recommended [17]. Conversely, foundation models designed for zero-shot use can organize developmental lineages and differentiation hierarchies without task-specific training [18], indicating that pretrained representations can encode continuous differentiation structure. Benchmark-driven evaluation has become an established genre in this journal, spanning single-cell foundation models [19, 20] and clinical question answering [21], and representation-quality assessment is emerging as a distinct concern in embedding evaluation [22]. Whether general-purpose representations adequately capture the ordered structure of a specific differentiation hierarchy, such as the corneal continuum, remains open; conversely, task-specific fine-tuning may improve discriminative accuracy while inadvertently collapsing the continuous geometric relationships among adjacent differentiation states that are biologically essential.

The fine-tuning literature has focused on whether adaptation improves classification accuracy and whether it causes forgetting of learned cell-type knowledge; full fine-tuning induces catastrophic forgetting at the label level, motivating parameter-efficient remedies [23]; a complementary line reports bounded fine-tuning gains and proposes marker-gene-guided post-pretraining instead [24]. With respect to geometry, a zero-shot benchmark reported geometric collapse of foundation-model embeddings during cellular dynamics reconstruction, with HVG/PCA outperforming them for trajectory recovery [25], and EMP-FM showed that naive fine-tuning often fails to resolve intermediate cell states under a dual classification–pseudotime objective [26]. Layer-resolved dissections of single-cell foundation models have begun to probe where biological information resides within the network [27, 28], and fine-tuned representations are already deployed in production annotation protocols without quantifying whether adaptation preserves ordinal structure [29]; a related cautionary literature documents cases in which well-performing models fail in settings that matter [30, 31]. What remains absent is a quantified measurement, under controlled fine-tuning conditions, of whether supervised adaptation preserves or destroys the continuous differentiation-state gradient that orders a hierarchy, which we term the *continuity problem*. The corneal epithelium is an ideal test system: its four ordered stages have well-established markers, so the expected ordinal structure is known *a priori* and gradient preservation can be measured directly in embedding

space. We evaluate a family of candidate gradient-preserving fine-tuning recipes under an identical locked protocol, including a continuity-regularized objective (CAFT), a cumulative ordinal head (CORN), an axis-anchored objective (AA-FT), a soft-Spearman rank surrogate, and a flow-consistency objective; whether any of them preserves ordered geometry at matched classification accuracy is a hypothesis to be tested on externally held-out, source-disjoint data, not an assumption.

Here we quantify the effect of classification fine-tuning on differentiation-state geometry in the normal human corneal epithelium. Using a donor-disjoint fair-split protocol on a 10-study corneal and ocular surface atlas (99,606 training cells from 49 donors; 20,394 test cells from 25 donors; zero donor overlap), we compare two reference arms (HVG/PCA with a linear classifier; frozen Geneformer-V2-104M with a linear probe) with adapted configurations sharing identical preprocessing, data splits, classification head, and training budget: naive cross-entropy fine-tuning, a continuity-regularized objective (CAFT), an ordinal head (CORN, with and without the continuity term), and three further gradient-preserving objectives (AA-FT, soft-Spearman, flow), evaluated on a same-corpus donor-disjoint holdout (GSE153515; 17,587 cells, 2 donors), a truly out-of-corpus external test (GSE299152; 21,374 epithelial cells, 3 donors), and a cross-tissue epidermal test set. Our contributions are fourfold: (1) a quantified, three-domain measurement of fine-tuning-induced differentiation-gradient collapse: catastrophic on same-corpus data (the *topology-preserving ordinal-axis collapse*, TP-OAC, layer-resolved across all thirteen hidden states), mild on truly out-of-corpus data, and reversed on cross-tissue data; (2) an ordinal head reparameterization (CORN) as the only recipe preserving the gradient at matched classification accuracy, and only within the pretraining corpus; (3) a systematic evaluation of five alternative gradient-preserving objectives (three honest negatives, one partial, one unstable), in the spirit of recent critical-evaluation studies in this journal [31, 32]; and (4) a contamination-aware scFM benchmark for the normal human corneal epithelium, pairing a same-corpus donor-disjoint holdout with a truly out-of-corpus external test under a locked multi-seed protocol with BH-FDR control and permutation nulls. These measurements define a collapse mode we term *topology-preserving ordinal-axis collapse* (TP-OAC), in which the ordinal differentiation axis is erased while discrete layer topology survives; it is distinct from the ordinal geometry that ordinal supervision induces [33], from pretrained layer profiles characterized without adaptation [34], and from the trajectory-fidelity trade-offs that trajectory benchmarks quantify across embedding methods [35]. Unlike prior evaluations of embedding quality, which focused on visualization stability [22] or zero-shot geometry [25], our measurements quantify how supervised adaptation itself alters ordinal differentiation geometry under controlled fine-tuning conditions. Our scope is restricted to normal human corneal epithelium; evaluation of disease states, including dry eye disease, is deferred to future work.

2. Results

2.1. Fine-tuning collapses the differentiation gradient while preserving layer topology

Classification fine-tuning collapsed the differentiation gradient of the corneal epithelium while leaving layer topology intact. Two non-adapted references (HVG/PCA with a linear classifier; frozen Geneformer-V2-104M with a linear probe) were compared against end-to-end cross-entropy fine-tuning and a continuity-regularized variant (CAFT), all sharing identical

preprocessing, data splits, classification head, and training budget, on a same-corpus, donor-disjoint holdout (GSE153515; Figure 1; three independent seeds per adapted arm; per-seed values in Table S1).

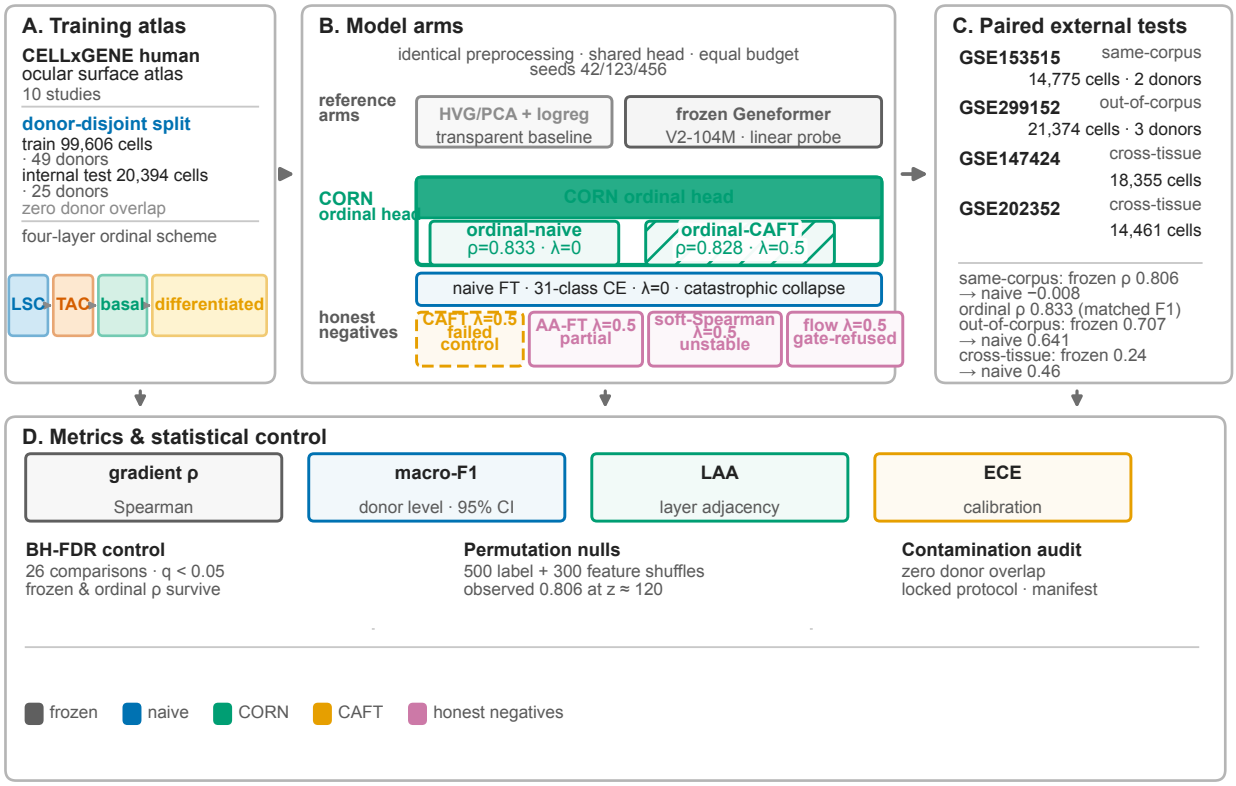


Figure 1: Study pipeline. Donor-disjoint training atlas (CELLxGENE), locked evaluation protocol, softmax-head and ordinal-head fine-tuning arms, and paired same-corpus (GSE153515) and post-cutoff out-of-corpus (GSE299152) external tests.

The two non-adapted reference arms preserved strong differentiation-gradient rank correlation on the same-corpus holdout: Spearman ρ was 0.8206 for HVG/PCA and 0.8064 for frozen Geneformer (Table 2). By contrast, naive supervised fine-tuning collapsed this ordinal signal to near zero: per-seed ρ ranged from -0.042 to $+0.054$ across seeds (three-seed mean -0.0082), with bootstrap 95% confidence intervals spanning zero in every run (Figure 2, Table 2). The sign of ρ was inconsistent across seeds: the residual ordinal correlation after fine-tuning was indistinguishable from noise. A randomized-head ablation – replacing the trained classifier with an untrained head before re-embedding – leaves the hidden states unchanged, ruling out that the collapse is a readout artifact of the trained classifier (Methods); the erasure resides in the fine-tuned body weights. The biological consequence of this collapse is quantified below (Figure 6).

Despite this gradient collapse, discrete layer topology was preserved. Layer adjacency accuracy (LAA) remained high in all four arms: 0.9771 (HVG/PCA), 0.9665 (frozen Geneformer), 0.9768 (naive, three-seed mean), and 0.9773 (CAFT, three-seed mean). Classification fine-tuning therefore disrupted the continuous ordering of the differentiation continuum while

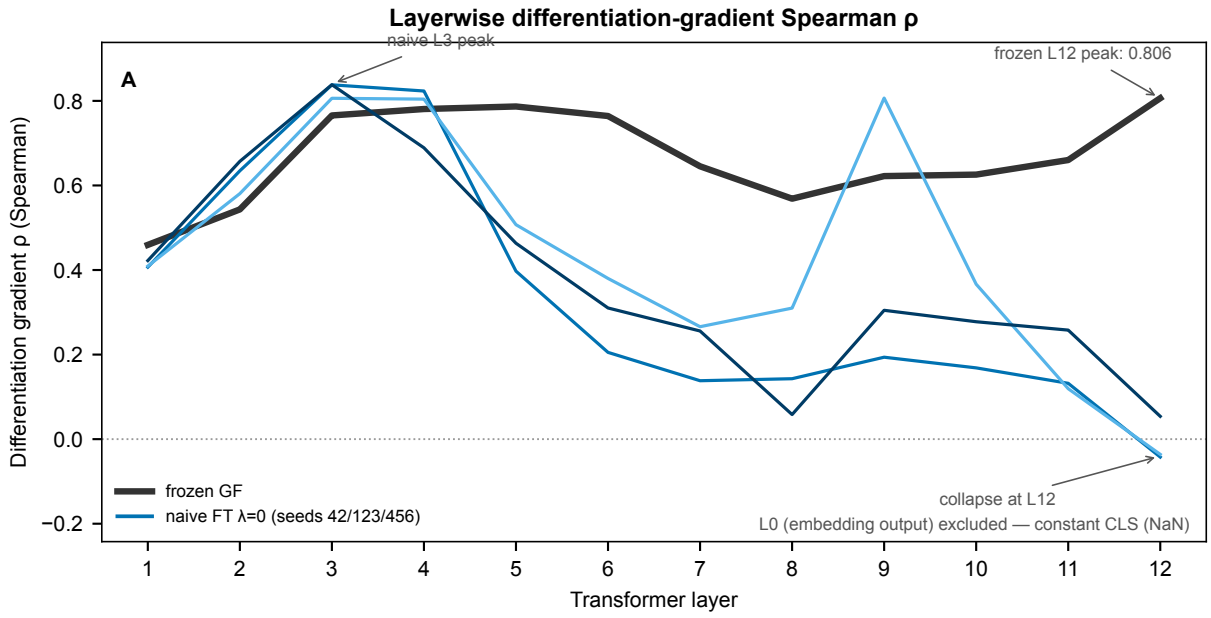


Figure 2: Layer-resolved analysis localizes the collapse to upper transformer layers. Layer-resolved differentiation-gradient Spearman ρ across hidden states L1–L12 for the Geneformer-based configurations (n = 14,775 external-test epithelial cells; three seeds per adapted arm; L0 excluded: constant [CLS]). Color key: thick dark gray line, frozen Geneformer; blue family, naive fine-tuning (seeds 42/123/456). The collapse is head-proximal: early layers retain the gradient, while L5–L12 decay. See also Figure S7 and Table S5.

leaving local layer-adjacency structure essentially intact: continuous gradient geometry and discrete layer topology are dissociable under task adaptation.

The continuity-regularized arm must be read as a failed control. A post hoc audit of the archived training code established that the continuity term never executed in any run, because of an implementation defect in the training hook, so those checkpoints trained pure cross-entropy, identical in objective to the naive arm (Methods, Provenance and contamination audit). The +0.210 separation in final-layer ρ therefore reflects run-to-run nondeterminism rather than a causal effect of the regularization, and classification performance was likewise indistinguishable: macro-F1 was 0.6451 (95% CI 0.6164–0.6738) for the continuity-regularized group versus 0.6492 (95% CI 0.6233–0.6751) for naive fine-tuning ($\Delta = -0.0041$; overlapping intervals). Throughout, we report this arm as an honest negative implementation result rather than a promoted safeguard. Fine-tuning nonetheless achieved the highest classification accuracy while collapsing the gradient to noise: differentiation geometry and classification performance are separable.

Layer-resolved analysis localized the collapse to upper transformer layers. To determine whether fine-tuning erases the differentiation-gradient signal globally or damages it in a layer-localized manner, we re-measured ρ at all thirteen hidden states (embedding output L0 and layers L1–L12; L0 excluded because the position-0 [CLS] representation is constant across cells) for all seven Geneformer-based configurations on the 14,775 external-test epithelial cells; the layer-resolved pipeline reproduced each model's single-value measurement exactly (Methods, Layer-resolved gradient measurement).

The collapse was concentrated in head-proximal layers, with the gradient intact in early

layers (Figure 2). First, early-layer representations preserved the gradient in every arm: L1–L4 ρ lay between 0.41 and 0.84, and the fine-tuned models reached their maximum at L3 (0.81–0.84 across seeds), matching or exceeding the frozen model at the same depth (L3, 0.7656). Second, the collapse observed at the final layer of the naive models was concentrated in the upper half of the network: from L5 (0.3973, seed 42) ρ decayed toward zero or negative values at L12, whereas the frozen model's gradient signal rose to its network-wide maximum of 0.8064 precisely at L12. Third, the continuity-regularized checkpoints ended at L12 ρ of 0.19–0.23, above the naive group's near-zero values but far below the frozen model's 0.8064; as the continuity term never executed (Methods), this difference is not attributable to the regularization objective.

Bootstrap confidence intervals at L12 (n = 1,000, cell-level resampling) quantified the separation: frozen $\rho = 0.8064$ [0.7918, 0.8164]; all naive seeds spanned zero; and the continuity-regularized checkpoints were positive (0.19–0.23) and, in this sample, separated from both groups. Because the continuity term never executed in that arm, the separation is a property of these particular cross-entropy runs, not evidence of a regularization effect. Gradient collapse under classification fine-tuning is therefore a head-proximal overwrite: the signal persists in intermediate layers and is selectively disrupted in the layers nearest the classification head.

To test whether this overwrite reflects the training process rather than an artifact of the trained classification head, we replaced the trained head of the naive seed-42 checkpoint with a freshly initialized, untrained random head: the layerwise differentiation-gradient ρ profile remained bit-identical at all thirteen layers (L12 $\rho = -0.0424$), while the same treatment on the frozen body preserved the frozen profile exactly (L12 $\rho = 0.8064$). The L5–L12 collapse is therefore intrinsic to the fine-tuned body weights, a training-induced rewrite of head-proximal layers, and not a readout artifact of the trained head (Methods, Layer-resolved gradient measurement).

Trajectory inference confirmed the geometric dissociation at the level of inferred cell ordering (Supplementary Figure S12). Diffusion pseudotime on the same external cells, rooted at the LSC class, correlated strongly with the reference layer-ordering pseudotime for the frozen model ($\rho = 0.850$, 95% CI 0.842–0.858) but weakly for naive fine-tuning ($\rho = 0.323$–0.348), with CAFT intermediate ($\rho = 0.372$–0.428). The loss of ordinal information under naive fine-tuning is thus observable not only in our gradient metric but in a widely used trajectory-reconstruction procedure.

2.2. The collapse is domain-specific: a three-domain boundary map

Having measured the collapse on a same-corpus holdout, we asked whether it reflects the evaluation domain rather than fine-tuning itself, and delineated its boundary across three domains of increasing distance from the pre-training corpus. To separate the effect of fine-tuning from the effect of pre-training exposure, we paired two external test sets with opposite corpus status and added a cross-tissue control: GSE153515 (normal limbal epithelium), whose primary publication DOI resolves to a row in the official Genecorpus-104M source table (53,290 cells, organ = eye) and thus lies within the pre-training corpus (least conservative); GSE299152, a single-cell map of the ocular surface glycocalyx [36], deposited after corpus assembly and absent from the table (post-cutoff, most conservative); and human epidermis (GSE147424 [37]) as a cross-tissue control. Construction and annotation of these sets follow the locked protocol in Methods (Supplementary Figure S13).

On the same-corpus holdout, fine-tuning drove performance while collapsing the gradient. On GSE153515, the HVG/PCA reference arm achieved donor-level macro-F1 of 0.6192 (95% CI 0.5747–0.6637), outperforming frozen Geneformer (0.4421; 95% CI 0.3749–0.5093), the opposite of the internal training-split ranking (0.8834 versus 0.7787). Frozen Geneformer showed poor calibration on external data (ECE 0.3010) and basal-class collapse (recall 0.216), whereas HVG/PCA remained well calibrated (ECE 0.0145) with strong gradient correlation. Both fine-tuned arms (naive 0.6492; CAFT 0.6451) outperformed both baselines. These results suggest that fine-tuning, not pretraining alone, drives cross-dataset generalization (Figure 3); the magnitude of these scores sits in the range expected for donor-disjoint validation, in which external macro-F1 falls well below internal-split numbers across the single-cell foundation-model literature [15, 16]. On this same-corpus holdout, the reference arms preserved the differentiation gradient (frozen $\rho = 0.8064$), whereas naive fine-tuning collapsed it to -0.008 with preserved layer topology (TP-OAC, above). The two evaluation dimensions are thus orthogonal at the level of individual arms: frozen Geneformer is best at geometry ($\rho = 0.8064$) but poorest at classification (macro-F1 0.4421), whereas the fine-tuned arms are best at classification (0.6492) but collapse the geometry; classification accuracy and gradient fidelity must be audited separately.

Both adapted arms generalized across the two donors (donor_1 > donor_2 in all six runs; Table 2). Because the TAC class is extremely rare (171/14,775 cells; 1.2%) and received zero true-positive predictions in every adapted run, we performed a sensitivity analysis excluding TAC (Supplementary Figure S10). Three-class macro-F1 rose by +0.21–0.22 in all adapted arms (naive 0.8599–0.8763; CAFT 0.8479–0.8786) and by +0.15 for frozen Geneformer (0.5921); the three-class confidence intervals did not overlap the four-class intervals, confirming that TAC rarity systematically depresses the primary endpoint. The naive-versus-CAFT comparison remained non-significant under the three-class definition, and the arm ranking was unchanged.

Out-of-corpus and cross-tissue domains bounded the collapse (Table 1, Figure 4). On the truly out-of-corpus test (GSE299152), classification fine-tuning collapsed the gradient only mildly: frozen $\rho = 0.707$ fell to 0.641 under naive fine-tuning (three-seed mean 0.641 ± 0.030), a partial rather than catastrophic loss, with a stable sign across seeds. The ordinal recipe did not transfer: ordinal-head ρ ranged from 0.35 (ordinal-naive; mean 0.350 ± 0.099) to 0.54 (ordinal-CAFT; mean 0.539 ± 0.007), below the naive arm's 0.641, and macro-F1 fell correspondingly (0.29–0.30 versus 0.452). The CORN head therefore preserves the gradient only within the pre-training corpus, where the ordinal axis is already encoded; out-of-corpus, its threshold predictions are poorly calibrated (ECE 0.56–0.61) and underperform cross-entropy.

On the cross-tissue domain (human epidermis [37], GSE147424; four ordered keratinocyte layers), the direction of the effect reversed: fine-tuning *improved* the gradient, from frozen $\rho = 0.24$ to $\rho = 0.46$ under naive fine-tuning (ordinal head 0.50). Notably, this inversion cannot be attributed to corpus absence: the DOI-level audit (Methods, Provenance and contamination audit) resolves GSE147424 to a row in the Genecorpus-104M source table (organ = skin, 45,200 cells), so the pretrained model did see epidermis, yet encoded its differentiation axis only weakly ($\rho = 0.24$ versus 0.81 for the cornea). The discriminating factor is therefore the strength of the tissue-specific differentiation axis that pretraining encodes, not corpus membership per se: classification supervision overwrites an axis it can already read (cornea)

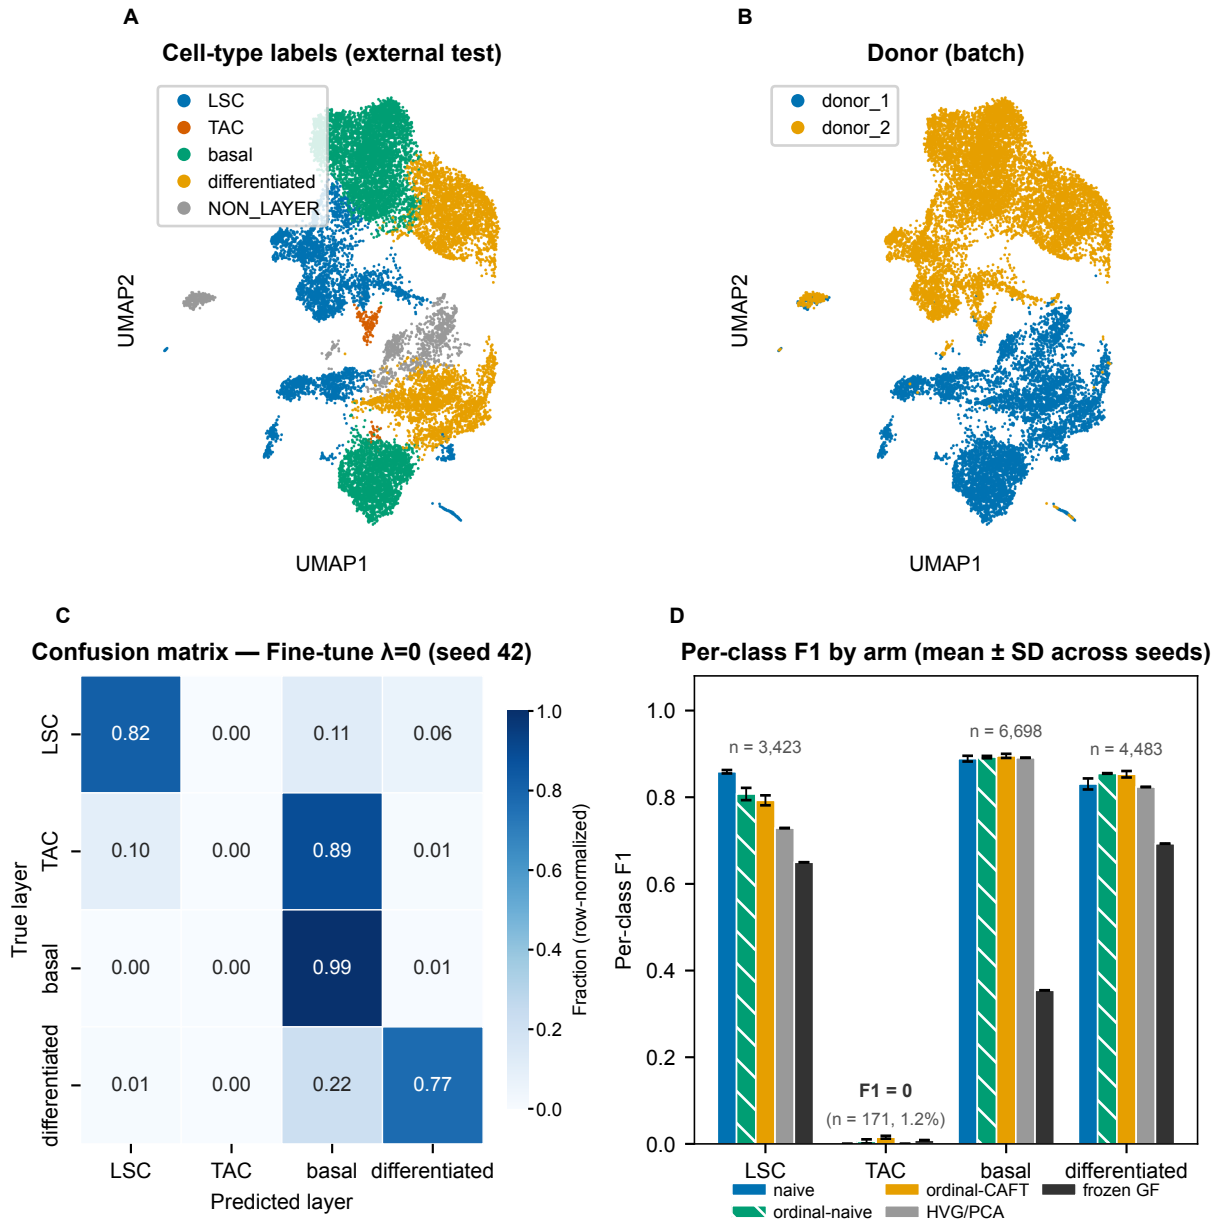


Figure 3: Same-corpus holdout: fine-tuning drives performance and collapses the gradient. External evaluation of adapted models on GSE153515 (n = 14,775 epithelial cells, two donors). (A) Cell-type UMAP. (B) Donor UMAP. (C) Confusion matrix. (D) Per-class F1 across five model arms (adapted arms three-seed means; reference arms single runs); the TAC column shows F1 = 0 (n = 171; 1.2%; three-class sensitivity in Figure S10). Color key (D): blue, naive; blue-green, ordinal-naive (hatched); orange, ordinal-CAFT; gray, HVG/PCA; dark gray, frozen. See also Figure S6.

and builds one where pretraining left it weak (epidermis). Together, the three domains delineate the boundary of TP-OAC: catastrophic on the in-corpus tissue with a strongly encoded axis, mild on post-cutoff out-of-corpus data, and absent (inverted) on the in-corpus tissue whose axis pretraining encodes weakly.

Table 1: Differentiation-gradient Spearman ρ (frozen, naive cross-entropy fine-tuning, and ordinal head) across three evaluation domains. Same-corpus adapted-arm values are three-seed means with per-seed ranges in parentheses; out-of-corpus adapted-arm values are three-seed mean ± SD, with the ordinal range spanning the ordinal-naive and ordinal-CAFT arm means; frozen and cross-tissue values are single runs. Statistical control (BH-FDR and permutation nulls) applies to the same-corpus family.

Domain	Frozen	CE-naive	Ordinal (COI
Same-corpus (GSE153515)	0.806	−0.008 (−0.042 to +0.054)	+0.833 (0.830–0.8
Out-of-corpus (GSE299152)	0.707	0.641 ± 0.030	0.350 ± 0.099 to 0.539 ± 0.
Cross-tissue epidermis (GSE147424)	0.240	0.460	0.
Cross-tissue epidermis (GSE202352)	0.520	0.710	0.

All primary same-corpus claims survived statistical control: a single BH-FDR family of 26 comparisons (q threshold 0.05) and permutation nulls (Methods, Statistical analysis plan). Frozen and ordinal ρ were significantly positive (both $q = 0.0086$), the naive collapse was not ($q = 1.0$), and the frozen observation sat more than 120 standard deviations from the label-shuffle null.

2.3. Alternative gradient-preserving recipes: honest negatives, a partial fix, and an unstable one

Is the collapse avoidable? Beyond the ordinal head, we evaluated five alternative gradient-preserving recipes under the identical locked protocol (Methods, Model architectures): three failed outright, one preserved the gradient only partially, and one was unstable across seeds (per-arm three-seed values in Table S9). CAFT, whose continuity objective never executed because of an implementation defect, is an honest negative: its higher final-layer ρ (+0.210 versus −0.008) was run-to-run nondeterminism rather than restoration (Methods, Provenance and contamination audit). A correctly wired continuity term (L_{cont}, tested inside the ordinal-CAFT arm) conferred no measurable benefit beyond the ordinal head ($\Delta\rho = 0.005$), echoing the auxiliary-signal suppression observed in frozen models [38]. The axis-anchored objective preserved the gradient only partially (three-seed mean 0.208), and the rank-based objective was unstable across seeds (mean 0.095, SD 0.172); flow was refused extension by its stability gate. Single-axis recipes cannot substitute for the ordinal head: a fixed axis captures at most the centroid-trajectory direction, not the full layer geometry that the metric derives per checkpoint (Methods, Axis definition revision).

2.4. An ordinal head parameterization preserves the differentiation gradient

The honest negative leaves open whether any fine-tuning recipe can preserve the differentiation gradient at matched accuracy. We therefore tested a constructive alternative, reparameterizing the classification head along the known ordinal axis instead of treating the four layers as unordered classes. Following the cumulative ordinal regression (CORN)

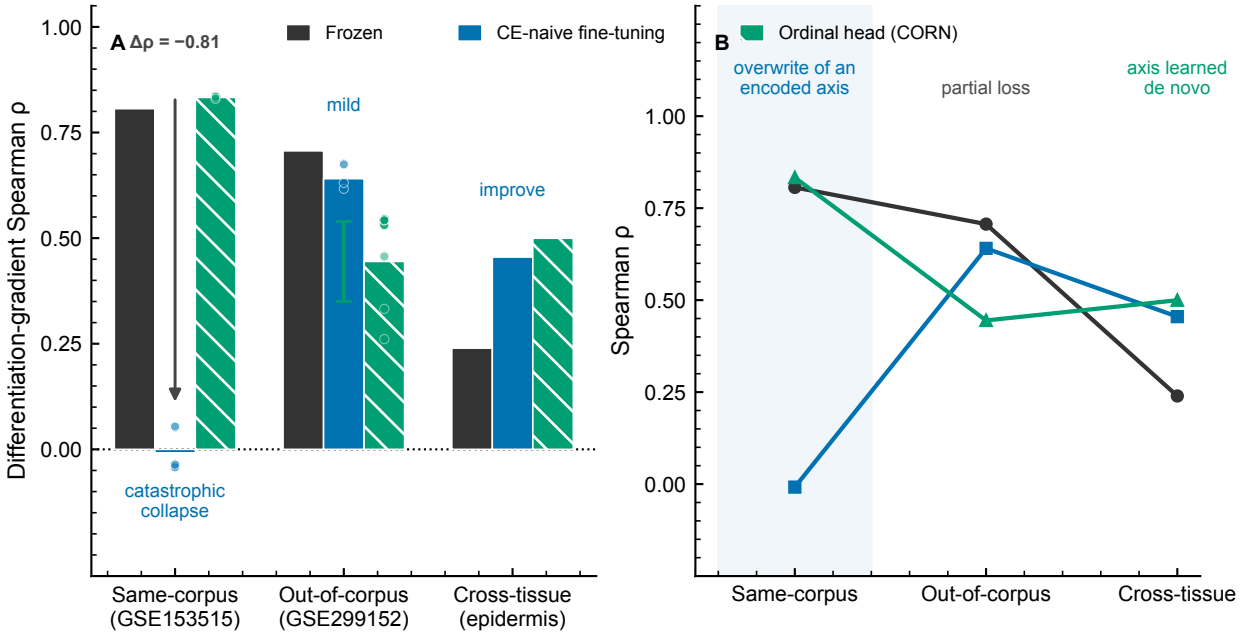


Figure 4: The collapse is domain-specific: a three-domain boundary map. (A) Spearman ρ by evaluation domain and arm (frozen, CE-naive, ordinal head); same-corpus adapted values are three-seed means, out-of-corpus ordinal range as mean ± SD (n = 14,775 same-corpus and 21,374 out-of-corpus epithelial cells; counts per domain in Methods). (B) Frozen-to-naive trajectories across domains: catastrophic overwrite within-corpus, partial loss out-of-corpus, and de novo axis learning cross-tissue. Color key: dark gray, frozen; blue, CE-naive; green, ordinal head (hatched). See also Table S5.

formulation [39], the head emits $K-1=3$ binary threshold logits, where threshold k predicts whether the cell has reached ordinal stage k or beyond; the loss is the sum of the three binary cross-entropies (masked to epithelial cells), and inference recovers the predicted layer as the number of thresholds with sigmoid activation ≥ 0.5. The CORN head makes no equidistance assumption on the thresholds and yields class probabilities directly for calibration ($P(y=k)=p_k-p_{k+1}$).

Two ordinal arms were trained under the identical locked protocol and budget: ordinal-naive ($\lambda=0$; ordinal loss only) and ordinal-CAFT ($\lambda=0.5$; ordinal loss plus the CAFT continuity term, here wired through a registered Trainer callback and therefore genuinely executed, in contrast to the archived CAFT arm). Both preserved the differentiation gradient at classification accuracy matched to the softmax arms (Table 2): three-seed mean ρ was 0.8330 (per-seed 0.8299–0.8351) for ordinal-naive and 0.8277 (0.8242–0.8321) for ordinal-CAFT, versus -0.008 for naive softmax fine-tuning, while donor-level macro-F1 was 0.6490 (95% CI 0.6100–0.6879) and 0.6479 (95% CI 0.6099–0.6860), statistically indistinguishable from the naive arm's 0.6492 (bootstrap 95% CI overlap). Layer adjacency accuracy remained $\approx$0.97 in both ordinal arms.

The near-equality of the two ordinal arms ($\Delta\rho=0.005$) indicates that the gradient preservation reflects the ordinal head parameterization itself rather than the continuity term: encoding the a priori layer order as ordered thresholds is sufficient, and the correctly wired centroid-adjacency regularization adds no measurable benefit on this axis. This is the constructive counterpart to the CAFT honest negative: a head reparameterization that respects the known ordinal structure preserves gradient geometry at matched classification performance, whereas the unordered softmax head erases it (Figure 5). This constructive result is confined to the same-corpus domain: on the truly out-of-corpus test the ordinal head preserved less gradient than cross-entropy fine-tuning (Table 1), indicating that the recipe operates by conserving an ordinal axis that pre-training has already encoded. The geometric preservation also comes at a calibration cost: on the same-corpus holdout, ordinal-naive ECE (0.1251) and ordinal-CAFT ECE (0.0609) were not better than the naive arm's 0.0575, and out-of-corpus the ordinal thresholds were markedly worse calibrated (ECE 0.56–0.61). The CORN head therefore trades calibration for geometry: it is a geometry-faithful readout whose probability estimates should be interpreted cautiously, particularly outside the pre-training corpus. The learned decision boundaries are themselves ordered, not merely accurate (Figure S11): the three threshold logits shift monotonically across differentiation stages, per-stage predicted probabilities align with the a priori order, and the per-cell cumulative monotonicity $P(y\geq 1)\geq P(y\geq 2)\geq P(y\geq 3)$ holds in 0.72 of cells versus 0.18 expected from an unordered head (three-seed mean 0.79; Supplementary Figure S11). Although the TAC class received no true-positive predictions in any adapted arm (F1 = 0), the ordinal head does not discard it: CORN threshold probabilities place TAC cells between LSC and basal (median predicted basal probability 0.975; Figure S11B), the biologically expected ordinal position, so the head's response to class rarity preserves rather than erases ordinal structure.

2.5. The differentiation axis recovers CNS-published biology

The geometric measurements above leave open whether the ordinal axis carries the biology established for this tissue, or whether the collapse is a rearrangement without biological content. We therefore asked whether each arm's axis recovers two *Cell*-published discoveries:

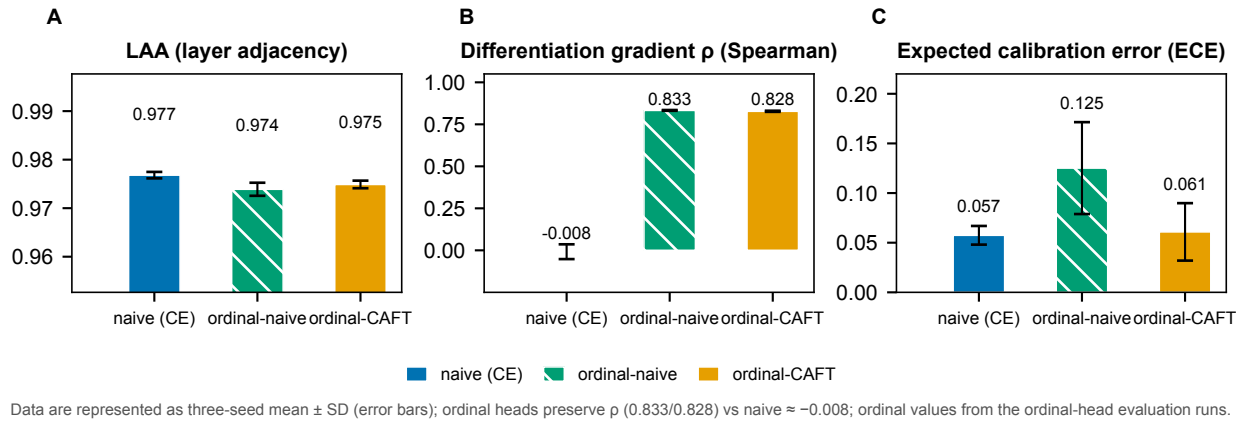


Figure 5: An ordinal head parameterization preserves the differentiation gradient. Continuity metrics under task adaptation. (A) Layer adjacency accuracy (LAA). (B) Differentiation-gradient Spearman ρ. (C) Expected calibration error (ECE). Bars denote the three adapted arms: naive (blue), ordinal-naive (blue-green, hatched), and ordinal-CAFT (orange) (three-seed mean ± SD; n = 14,775 cells per arm). See also Figure S3.

Table 2: External test performance by model arm (GSE153515, four classes; n = 14,775 epithelial cells from two donors). Adapted-arm values are three-seed means with donor-level bootstrap 95% CIs; reference arms are single runs (seed 42). LAA, layer adjacency accuracy; ECE, expected calibration error.

Model arm	Macro-F1	95% CI	LAA	ρ	ECE
HVG/PCA + linear	0.6192	[0.5747, 0.6637]	0.9771	0.8206	0.0145
Frozen Geneformer	0.4421	[0.3749, 0.5093]	0.9665	0.8064	0.3010
Naive FT ($\lambda = 0$)	0.6492	[0.6233, 0.6751]	0.9768	−0.0082	0.0575
CAFT ($\lambda = 0.5$)	0.6451	[0.6164, 0.6738]	0.9773	+0.2102	0.0600
Ordinal-naive ($\lambda = 0$)	0.6490	[0.6100, 0.6879]	0.9739	+0.8330	0.1251
Ordinal-CAFT ($\lambda = 0.5$)	0.6479	[0.6099, 0.6860]	0.9749	+0.8277	0.0609

the slow-cycling limbal stem-cell compartment of Cotsarelis et al. [1] (a quiescent limbal stem-cell population that can be stimulated to proliferate, yielding transient-amplifying cells and then terminally differentiated keratinocytes) and the basal-to-suprabasal keratin switch of Fuchs and Green [40]. Along the frozen axis the published program is recovered (Figure 6): the stemness program (KRT15, KRT19, TP63) declines monotonically from the stem-cell pole toward differentiation ($\rho = -0.41$; bootstrap 95% CI -0.42 to -0.39), the corneal differentiation program (KRT12, KRT3) rises ($\rho = +0.69$; $+0.68$ to $+0.70$), and the cycling program peaks in the transient-amplifying compartment (mean score 3.15 in TAC cells versus < 0.05 in every other class; PCNA and MCM2, both absent from the annotation marker sets, peak in the same compartment). Naive fine-tuning erases the graded structure: the stemness correlation collapses to a three-seed mean of $+0.02 \pm 0.03$ (per-gene: KRT15 $-0.53 \to -0.06$; KRT19 $-0.41 \to -0.12$), while the differentiation-pole signal is only partially retained ($+0.69 \to +0.29 \pm 0.03$), consistent with retained coarse class discrimination (macro-F1 0.65) without the graded axis. The ordinal head, which preserves the gradient geometry, also preserves the biology with negligible seed variance: three-seed stemness ρ was -0.427 ± 0.004 (ordinal-naive) and -0.407 ± 0.012 (ordinal-CAFT), and differentiation ρ $+0.557 \pm 0.005$ and $+0.567 \pm 0.012$, statistically indistinguishable from the frozen reference.

The recovery is not an annotation artifact: the per-gene panel includes markers absent from every locked annotation marker set (NFIB -0.16, CDH3 -0.07, PCNA $+0.14$ on the frozen axis; NFIB -0.22 on the ordinal axis), and all program correlations exceed the permutation null by tens of standard deviations (stemness $z = -49$; differentiation $z = +82$); TP63 alone is basal-peaked rather than stem-cell-peaked in this 10x dataset and is not axis-separable ($+0.13$), an honest boundary of single-marker recovery. The same framework transfers to the cross-tissue control, and the reversal is reproduced on a second independent epidermal dataset (GSE202352 [41]; Figure 6C). On GSE147424 the basal keratin program (KRT5, KRT14, COL17A1) declines along the axis in every arm (-0.59 frozen to -0.75 fine-tuned), and the cycling program declines from the proliferative basal compartment (-0.23 fine-tuned), recovering the Fuchs–Green switch in its measurable (basal-off) half; the differentiation-pole keratins are only partially monotone in this dataset (KRT10 $+0.18$ ordinal; KRT1 bimodal across layers), a limitation we report rather than smooth over. On GSE202352 the frozen axis instead encodes the differentiation-pole keratins (KRT1, KRT10; $\rho \approx +0.83$) rather than the basal pole ($+0.11$), but every adapted arm builds the basal-off program (-0.65 to -0.79) with the cycling program declining from the proliferative basal compartment (-0.21), and the class-ordinal reversal holds (frozen 0.52 $\to$ fine-tuned 0.71–0.76). The frozen baselines thus differ across datasets – the axis encodes different halves of the switch – while adapted arms consistently orient it toward the full basal-to-suprabasal program, and naive fine-tuning erases the graded biology on the in-corpus cornea in lockstep with the geometric collapse. Together with the cross-tissue reversal, these data indicate that the axis is the biology: the frozen and ordinal embeddings arrange cells along the differentiation programs discovered in *Cell* 1980 and 1989, and naive fine-tuning erases that graded biological structure while leaving layer topology intact.

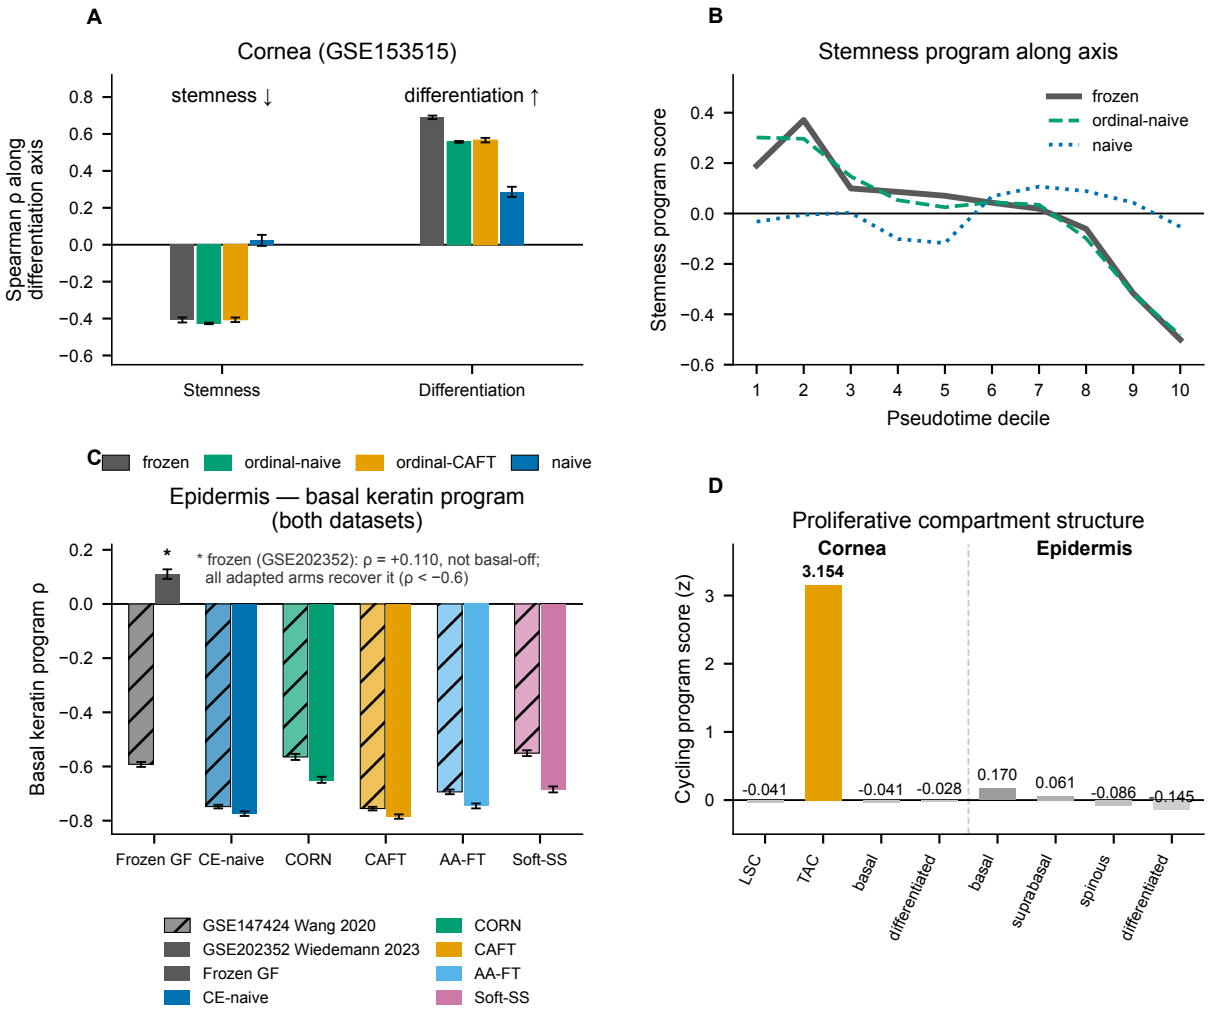


Figure 6: The differentiation axis recovers CNS-published biology. (A) Cornea (GSE153515, n = 14,775): Spearman ρ of stemness and differentiation programs along each arm's axis (frozen with bootstrap 95% CI; adapted arms three-seed mean ± SD). (B) Stemness-program score by pseudotime decile: the graded decline is preserved by frozen and ordinal embeddings and flattened by naive fine-tuning. (C) Epidermis, both datasets (GSE147424, n = 18,355; GSE202352, n = 14,461): basal-keratin program ρ by arm; every adapted arm recovers the basal-off switch, reproducing the cross-tissue reversal. (D) Cycling program by class: the proliferative transient-amplifying compartment (corneal TAC; epidermal basal) sits between the quiescent stem pole and the post-mitotic layers, the slow-cycling structure of Cotsarelis et al. [1]. Color key follows Figure 5. See also Methods, CNS-anchored marker-program validation.

3. Discussion

The central finding of this study is that differentiation geometry and classification performance are separable properties of single-cell representations. Classification fine-tuning of Geneformer-V2-104M collapsed the differentiation-gradient Spearman ρ from $\approx$0.82 (0.8206 for HVG/PCA; 0.8064 for frozen Geneformer) to -0.008 on the same-corpus holdout, while layer adjacency accuracy remained $\approx$0.97 in all adapted arms and macro-F1 rose to 0.6492. The layer-resolved analysis localizes the loss to a head-proximal overwrite: intermediate layers retain the gradient (L3 ρ 0.81–0.84, matching the frozen model at the same depth), while the readout layers nearest the classification head (L5–L12) are selectively degraded. Pretraining concentrates differentiation structure precisely in the final layer, and classification training rewrites exactly that layer. The practical implication is direct: pipelines that fine-tune foundation models and then reuse the resulting embeddings for analyses that require continuous structure (trajectory inference, pseudotime ordering, state-transition quantification) will lose ordinal information even when classification accuracy is high; where continuous structure matters, intermediate-layer or frozen readouts deserve explicit consideration. Conversely, the ordinal recipe that preserves geometry does so at a calibration price (same-corpus ECE 0.1251 versus 0.0575 for naive; out-of-corpus 0.56–0.61), so the trade-off among gradient fidelity, classification accuracy, and calibration is three-way rather than binary.

Three observations bound this collapse and its mechanism. First, the effect is domain-specific rather than universal: catastrophic on same-corpus data, mild on a truly out-of-corpus test (frozen 0.707 $\rightarrow$ naive 0.641), and reversed on cross-tissue data (frozen 0.24 $\rightarrow$ naive 0.46). The inversion is explained by the strength of the pretrained, tissue-specific differentiation axis rather than by corpus membership: the corpus audit places the epidermal test within Genecorpus-104M (GSE147424 resolves to an organ = skin row; 45,200 cells), yet the frozen corneal axis is strongly organized (ρ = 0.81) while the frozen epidermal axis is weak (ρ = 0.24); classification supervision overwrites an axis it can already read and builds one where pretraining left it weak. Second, the collapse is recipe-specific: of five gradient-preserving objectives evaluated under an identical protocol, only the ordinal head preserved the gradient robustly, and only within-corpus; the correctly wired continuity term conferred no benefit beyond the ordinal head itself, and single-axis recipes captured at most the centroid-trajectory direction. Together with the head-proximal overwrite, these observations indicate that preserving ordinal geometry under adaptation requires encoding the ordinal structure in the supervisory signal itself, at the readout where the collapse occurs.

We frame this pattern as *topology-preserving ordinal-axis collapse* (TP-OAC): supervised adaptation erases the ordinal axis of a pre-existing differentiation continuum while cluster membership and layer topology survive, with the loss graded by depth and concentrated in head-proximal layers. TP-OAC joins the neural-collapse family of phenomena [42–44] yet differs from each of them: neural collapse is the terminal-phase convergence of within-class features, whereas TP-OAC is a cross-class ordinal-axis erasure that leaves discrete topology intact. It is likewise distinct from the ordered geometry that ordinal supervision induces [33], from the intermediate-depth localization of differentiation signal under zero-shot conditions [34], from the generic representational degradation that R3F mitigates [45], and from the label-level catastrophic forgetting documented for parameter-efficient fine-tuning [23]; our randomized-head ablation localizes the effect to the fine-tuned body rather than the readout.

Notably, zero-shot evaluations found that Geneformer's embedding space fails to retain discrete cell-type information [46], yet we measured a strong frozen ordinal gradient ($\rho = 0.8064$): ordinal geometry and discrete type identity are dissociable properties of the same representation, and each must be audited separately. The frozen cross-architecture probes extend the architectural scope: scFoundation, UCE, and scGPT-human all separated the four layers but encoded the ordinal axis with markedly different strength ($\rho = 0.75$, 0.58, and 0.22 on GSE153515; Table S8), so ordinal-geometry encoding is architecture-dependent rather than universal. To our knowledge, fine-tuning-induced collapse of an existing ordinal axis, with topology survival, has not previously been measured.

The biological validation anchors these geometric measurements in published biology. Along the frozen and ordinal axes, the stemness program declines and the differentiation program rises monotonically, recovering the slow-cycling limbal stem-cell structure of Cotsarelis et al. [1]; on epidermis, every arm recovers the basal keratin program of Fuchs and Green [40]. Naive fine-tuning erases the graded biology in lockstep with the geometric collapse, so the erasure is a biological loss, not a mere rearrangement. Classification accuracy alone is therefore an insufficient criterion for evaluating adapted foundation models in differentiation biology: geometry-aware, layer-resolved metrics expose losses that aggregate accuracy conceals. Fine-tuning benefits are setting-dependent [15, 16]; adaptation can preserve gene–gene relationships when the context matches pretraining [47], so collapse reflects a mismatch between the supervisory signal and the geometry to be preserved, not an inevitable consequence of adaptation. Encoding biological structure into the supervisory signal is one productive route: hierarchical cross-entropy losses over cell-ontology graphs improve out-of-distribution annotation [48], and our ordinal head extends the same principle from classification generalization to the preservation of ordinal geometry; EMP-FM's positive dual-task result [26] is consistent, because pseudotime regression provides a direct ordinal target for every cell, whereas centroid-based adjacency constrains only the already-overwritten readout layer. The contamination-aware paired benchmark released with this study is intended as a validation resource for such geometry-aware evaluation, in corneal biology and beyond: evaluations in this field remain predominantly within single datasets and settings, and leakage-prone splits can inflate apparent performance [49]. Geometry-faithful evaluation of this kind is a prerequisite for the reliable virtual-cell representations envisioned by the AI Virtual Cell initiative [50] and complements multi-capability benchmarks such as VCBench [51]. Reporting the gate-refused and unstable recipe arms as honest negatives is deliberate: the pre-specified stability gate was applied without post hoc tuning, and publishing such outcomes, paired with a constructive layer-resolved diagnostic, serves the community better than selective reporting [31, 32].

3.1. Limitations of the study

This study has four principal limitations. First, it evaluates a single model family (Geneformer-V2-104M) on a single tissue (normal human corneal epithelium); the external evidence base comprises a two-donor same-corpus holdout, a three-donor post-cutoff out-of-corpus test, and two cross-tissue epidermal datasets (GSE147424 and GSE202352), and whether the dissociation between gradient geometry and layer topology, the head-proximal overwrite pattern, and the domain-specific boundary generalize across architectures, tissues,

and larger donor panels remains to be tested. The cross-architecture probes (frozen scFoundation, UCE, and scGPT-human) each separated the four layers (probe macro-F1 0.79–0.95) but ranged from strong (scFoundation $\rho = 0.75$) to weak (scGPT-human $\rho = 0.22$) ordinal encoding, with all external tests out-of-corpus for these models. The domain-specific collapse map reported here is therefore measured within Geneformer-V2-104M on normal human corneal epithelium; transfer across architectures and tissues is left to dedicated evaluation. In particular, a post-cutoff (out-of-corpus) single-cell epidermal test, which would separate the axis-encoding-strength effect from any residual corpus-exposure effect on the cross-tissue reversal, is left to future work. Second, the TAC class achieved F1 of 0.0 in every adapted run on GSE153515, reflecting extreme class rarity (171 of 14,775 external-test cells; 1.2%) rather than a method-specific failure, and the TAC class was absent altogether from GSE299152 (proliferation markers below 0.6% of cells); although our sensitivity analysis confirms that excluding TAC raises macro-F1 uniformly across arms without changing the ranking, conclusions regarding TAC representations are constrained by sample size in both test sets. Third, the gradient-preserving recipe arms beyond the ordinal head were expanded to three seeds only for the two non-gated arms (AA-FT $\lambda = 1.0$, soft-Spearman $\lambda = 0.5$); flow was refused extension by its stability gate. The three-seed results confirm that AA-FT preserves the gradient only partially and that soft-Spearman is unstable across seeds, so neither is claimed as a reliable fix; these results bound the search space for gradient-preserving adaptation, directing future recipe design toward intermediate-layer constraints. Fourth, the pre-training audit places GSE153515 within the corpus at the DOI level (53,290 cells, organ = eye; the V1 audit found zero corneal datasets), places the cross-tissue epidermal control GSE147424 within the corpus as well (45,200 cells, organ = skin), finds 10.4% of training-atlas cells overlapping corpus DOIs, and detected no cell-level overlap in bounded Signal A sampling (0.54% of the corpus). The frozen model's same-corpus performance may therefore be inflated by pre-training exposure; our conclusions rest on the relative, fine-tuning-induced measurements, which are insensitive to any shared pre-training contribution, and the cross-tissue reversal is interpreted as an axis-encoding-strength effect (the frozen epidermal axis is weak, $\rho = 0.24$, despite in-corpus skin exposure) rather than a corpus-absence effect.

4. Methods

4.1. Data sources and preprocessing

Training atlas. Training and internal-test data were derived from the "Single cell atlas of the human ocular surface" collection (CELLxGENE Discover, collection identifier 0f7d022a-46c7-4e64-be4c-e34adbb78089). The data file used is the current dataset-version of CELLxGENE dataset 8ef06ea4-8d0f-4303-a652-7d4604402b1f (schema 7.1.0, published 2025-05-05), distributed as file 30127f02-d0b9-4b56-bf4c-e5b77e4989bd.h5ad (10.29 GB; MD5 39f098367af95bc3e1ee92528ceb251e; the file identifier is the dataset-version identifier at time of download, retained for traceability). The atlas is a pre-publication consortium product (CZI Cell Science / Human Cell Atlas Eye Network) without an associated journal publication or DOI, and is therefore cited as a CELLxGENE collection [9]. The collection aggregated ten constituent studies spanning corneal, limbal, scleral, and fetal ocular surface tissues [3–6, 8, 52, 53]. The aggregated AnnData object contained 120,000 cells with 20 curated cell-type labels and 31 author-assigned cell-type labels.

Cell-type labels were harmonized into a five-layer ordinal scheme reflecting the corneal epithelial differentiation continuum: LSC (including LPC and fetal epithelial progenitors; layer index 0), TAC (limbal epithelium; index 1), basal (Epi-PMC; index 2), wing (Epi-TDC/PMC; index 3), and squamous (Epi-TDC/unknown epithelium; index 4). Non-epithelial populations (immune, endothelial, stromal, conjunctival, melanocyte) were assigned layer index -1 (NON_LAYER) and excluded from ordinal analyses. For the four-class classification task evaluated on external data, the scheme was collapsed to LSC, TAC, basal, and differentiated (wing and squamous merged).

External test datasets. The study paired two external test sets with opposite pre-training-corpus status (Methods, Provenance and contamination audit). GSE153515 (GEO accession), a 10x Genomics scRNA-seq dataset of normal human limbal basal epithelium [3, 6], contained 17,587 cells from two donors (GSM4646295, GSM4646296) with an initial gene matrix of 32,738 features and is the same-corpus holdout. GSE299152 [36], a single-cell map of the ocular surface glycocalyx in adult human and mouse (human cells used here), contained 26,628 human cells from three donors with an initial gene matrix of 36,601 features and is the post-cutoff, out-of-corpus test. Disjointness from the training atlas was verified programmatically via the CELLxGENE Discover API: no overlapping GEO accession, donor identifier, or dataset hash was found between either external test and any of the ten constituent training studies; disjointness from the pre-training corpus is addressed in the contamination audit.

4.2. Fair donor-disjoint splitting protocol

To prevent donor-level information leakage between training and evaluation, a strict donor-disjoint split was applied to the training atlas prior to any model fitting. Donors were partitioned using a deterministic random assignment (seed = 42), yielding 49 donors (99,606 cells) for training and 25 donors (20,394 cells) for internal testing, with zero donor overlap between partitions. The split protocol ensured representation of all 20 curated labels and all 31 author-assigned labels in both partitions. The split manifest, including per-partition cell counts, donor lists, label distributions, and SHA-256 checksums of the resulting train and test AnnData files, was locked and archived (artifacts/paper1_fair_split/manifest.json). No cells were reassigned after manifest generation.

4.3. External test annotation pipeline

Because GSE153515 was deposited with an empty observation metadata table, a fully reproducible de novo annotation pipeline was implemented (scripts/protocol_15_annotate_gse153515.py): (1) donor assignment from the 10x barcode suffix convention; (2) marker scoring on the full 32,738-gene matrix using eight curated marker gene sets spanning LSC, TAC, basal, differentiated, conjunctival, melanocyte, immune, and stromal identities; (3) quality-control filtering by minimum gene count and maximum mitochondrial fraction (15,827 of 17,587 cells retained, 89.9%, spanning 10,605 genes); (4) integration and clustering: highly variable gene selection, log-normalization, scaling, PCA, Harmony batch correction (n_pcs = 30) [54], and Leiden clustering (resolution 0.6) yielding 16 clusters; (5) DEG-based cell-type assignment via Wilcoxon rank-sum test with an enrichment score $(n_{\text{in top200}}/n_{\text{total}}) \times (1 - \text{mean_rank}/200)$ against the top 200 cluster-specific DEGs, with clusters below an epithelial enrichment threshold of 0.3 classified as NON_LAYER; and (6) harmonization

to the four-class scheme (basal 6,698; differentiated 4,483; LSC 3,423; TAC 171 cells) with non-epithelial cells excluded. All annotation parameters were locked in a configuration file (config/external_test_gse153515.yml) before any model evaluation. The second external test, GSE299152, was annotated with a mirrored pipeline (scripts/protocol_15b_annotate_gse29 9152.py; config/external_test_gse299152.yml) with two recorded deviations: Endothelial and Stromal marker sets were added to the NON_LAYER gate, and Leiden resolution was auto-tuned to 0.7 (15 clusters). After QC (23,970 of 26,628 cells retained, 90.0%; 11,274 genes), harmonization yielded LSC 15,921, basal 1,099, and differentiated 4,354 cells, with an empty TAC class (proliferation markers below 0.6% of cells) and 2,596 non-epithelial cells excluded; low-confidence clusters (melanocyte enrichment 0.199 below the 0.3 gate; MT-high and stress clusters with near-zero enrichment) were default-assigned to LSC/LPC and flagged in the annotation report.

4.4. Model architectures

Four model configurations were evaluated under an identical experimental protocol (same preprocessing, data split, classification head, training epochs, and random seeds); two additional ordinal-head arms (Arms 5–6) and five gradient-preserving objective arms (Arms 7–11) shared the same protocol and budget:

Arm 1—HVG/PCA + linear classifier (transparent baseline). Highly variable genes were selected from the training partition; expression values were log-normalized and scaled, PCA was applied, and a linear classifier was trained on the resulting principal components.

Arm 2—Frozen Geneformer + linear probe. Geneformer-V2-104M (BertForMaskedLM architecture; hidden dimension 768, 12 layers, 12 attention heads, vocabulary 20,275) was loaded from the Hugging Face model hub (snapshot commit 04c2b2e84da7c0f385c3f9ad8f3ec24bab6650e5; model weights SHA-256 fff5cba2...), pretrained on Genecorpus-104M ($\sim$104 million cells from 2,903 datasets) [11, 12]. All Transformer parameters were frozen; only a linear classification head appended to the [CLS] token representation was trained.

Arm 3—Geneformer naive fine-tuning ($\lambda = 0$). The full model was fine-tuned end-to-end with a standard cross-entropy classification loss over the 31 author-assigned cell-type classes (no continuity regularization), serving as the fair adaptation comparator; predictions were merged to the four epithelial differentiation classes for external evaluation.

Arm 4—Geneformer + CAFT ($\lambda = 0.5$). The full model was fine-tuned with the CAFT composite loss (Methods, CAFT loss formulation) at a single fixed regularization strength, testing CAFT as a candidate safeguard for ordered differentiation geometry.

Arm 5—Geneformer + ordinal-naive head ($\lambda = 0$). The full model was fine-tuned end-to-end with the CORN ordinal loss (Methods, Ordinal head formulation) and no continuity regularization.

Arm 6—Geneformer + ordinal-CAFT head ($\lambda = 0.5$). The full model was fine-tuned with the CORN ordinal loss plus the CAFT continuity term, wired through a registered Trainer callback so that the continuity objective was genuinely executed (in contrast to the archived CAFT arm).

Arms 7–11—axis-anchored, soft-Spearman, and flow objectives. Three further gradient-preserving objectives were evaluated with a four-class cross-entropy head on the

revised centroid-PCA differentiation axis (Methods, Axis definition revision): Arm 7 (AA-FT, $\lambda = 0.5$) and Arm 8 (AA-FT, $\lambda = 1.0$) augment the classification loss with a projection-constraint term pulling embeddings onto the fixed differentiation axis [55]; Arm 9 (soft-Spearman, $\lambda = 0.5$) and Arm 10 (soft-Spearman, $\lambda = 1.0$) add a differentiable Spearman surrogate (soft rank via sigmoid pairwise comparisons with a diagonal correction, ABAG-Rank formulation [56], with a margin surrogate [57]; a differentiable sorting-and-ranking method [58]); and Arm 11 (flow, $\lambda = 0.5$) adds a flow-consistency objective constraining embeddings toward ordered transitions between adjacent layers. These arms used the same preprocessing, data split, and training budget as Arms 3–6.

Arms 3–6 were each trained with three independent random seeds (42, 123, 456) for three epochs. Arms 7–11 were trained with seed 42 only: a three-seed expansion was gated on the seed-42 evaluation reaching a pre-specified differentiation-gradient range (0.5–0.7) on the same-corpus holdout, which no arm reached (Results). Training used a per-device batch size of 8 with gradient accumulation of 4 (effective batch size 32), bfloat16 mixed precision, gradient checkpointing, linear learning-rate warmup over 500 steps, weight decay 0.001, and learning rate 5×10^{-5}.

4.5. CAFT loss formulation

The CAFT objective augmented standard cross-entropy classification with a continuity regularization term:

$$\mathcal{L} = \mathcal{L}_{\text{cls}} + \lambda \cdot \mathcal{L}_{\text{cont}}, \tag{1}$$

where $\mathcal{L}_{\text{cls}}$ is the cross-entropy classification loss over the 31 author-assigned cell-type classes of the training atlas (merged to the four epithelial differentiation classes for external evaluation) and $\mathcal{L}_{\text{cont}}$ penalizes deviation from expected adjacency structure. For each training cell x_i with assigned differentiation layer $\ell_i \in \{0, 1, 2, 3, 4\}$, the continuity loss was

$$\mathcal{L}_{\text{cont}} = -\frac{1}{N} \sum_{i=1}^{N} \sum_{\ell' \in \mathcal{A}(\ell_i)} \cos\Big(\mathbf{z}_i,\ \text{sg}(\mathbf{c}_{\ell'})\Big), \tag{2}$$

where $\mathbf{z}_i$ is the ℓ_2-normalized [CLS] embedding of cell x_i, $\mathcal{A}(\ell_i)$ denotes the adjacent layers ($\ell_i \pm 1$, clipped to $[0, 4]$), $\mathbf{c}_{\ell'}$ is the centroid of training embeddings assigned to layer ℓ', and $\text{sg}(\cdot)$ is the stop-gradient operator. Layer centroids were precomputed at the start of each epoch from detached embeddings and held fixed throughout that epoch. NON_LAYER cells ($\ell = -1$) were excluded from the continuity term but contributed to $\mathcal{L}_{\text{cls}}$.

4.6. Ordinal head formulation

To test whether the collapse of the differentiation gradient is caused by treating the four ordered layers as unordered softmax classes, Arms 5–6 replaced the softmax classification head with a cumulative ordinal head following the CORN formulation [39]. For $K = 4$ ordinal classes, the head emits $K - 1 = 3$ binary threshold logits, where threshold k predicts whether the cell has reached ordinal stage k or beyond ($P(y \geq k)$); the loss is the sum of the three binary cross-entropies, masked to epithelial cells:

$$\mathcal{L}_{\text{ord}} = \sum_{k=1}^{K-1} \text{BCE}\Big(\sigma(f_k(\mathbf{z})),\ [y \geq k]\Big), \tag{3}$$

where $f_k(\mathbf{z})$ is the k-th threshold logit on the [CLS] representation $\mathbf{z}$, σ is the sigmoid, and $[y \geq k]$ is the binary target indicating that the cell's ordinal layer is at least k. At inference the predicted class is the number of thresholds with $\sigma(f_k(\mathbf{z})) \geq 0.5$. The CORN head makes no equidistance assumption on the thresholds and yields class probabilities directly for calibration ($P(y = k) = p_k - p_{k+1}$, with $p_0 = 1$, $p_K = 0$). The ordinal-CAFT arm (Arm 6) added the CAFT continuity term ($\mathcal{L} = \mathcal{L}_{\text{ord}} + \lambda \cdot \mathcal{L}_{\text{cont}}$), with $\mathcal{L}_{\text{cont}}$ wired through a registered `TrainerCallback` so that the continuity objective was genuinely executed at each epoch start, in contrast to the archived CAFT arm, where the equivalent hook was never dispatched.

4.7. Evaluation metrics

All metrics were computed on the locked external test sets using the Tier-B evaluation pipeline (corneal_fm/evaluation.py).

Primary endpoint. Donor-level macro-averaged F1 (macro-F1) on the four-class external test set, with 95% confidence intervals estimated by nonparametric bootstrap resampling of donors (n_boot = 1,000; seed = 42). Per-class precision, recall, F1, and support were reported for each arm and seed.

Secondary endpoints. (i) Differentiation-gradient rank correlation (Spearman ρ): an ordinal-structure metric computed solely from the layer-centroid geometry of the learned embeddings, without marker-gene expression. For each of the four layer classes (NON_LAYER excluded), the centroid of the cell embeddings was computed; the principal differentiation axis was defined as the first singular vector of the centered layer centroids, oriented so that pseudotime increases from LSC to differentiated; each cell was assigned a pseudotime equal to the projection of its embedding onto this axis; and ρ was the Spearman correlation between per-cell pseudotime and the a priori layer ordinal. (ii) Layer Adjacency Accuracy (LAA): the fraction of nearest-neighbor pairs in the learned embedding space belonging to adjacent or identical differentiation layers. (iii) Expected Calibration Error (ECE): the weighted mean absolute difference between predicted confidence and observed accuracy across probability bins. (iv) Internal silhouette coefficient and gap statistic were computed as diagnostic-only measures and were not used for model selection or promotion decisions.

Layer-resolved gradient measurement. At inference, all thirteen hidden states (embedding output L0 and layers L1–L12) were collected for each of the 14,775 external-test epithelial cells, and the differentiation-gradient ρ was recomputed at each depth from the position-0 [CLS] representation (scripts/measure_layerwise_rho.py). A sanity gate required the L12 ρ of every model to reproduce the corresponding single-value measurement within 10^{-6} (all seven models passed bit-exactly). At L0 the position-0 [CLS] representation is constant across cells, so ρ is undefined (NaN) and was excluded (expected behavior). For the randomized-head ablation, the trained classification head was replaced with a freshly initialized, untrained head and hidden states were collected without further training; the maximum hidden-state difference from the intact model was 0 for both the fine-tuned and frozen bodies. Cell-level bootstrap 95% confidence intervals at L12 were computed with n_boot = 1,000 resamples (seed = 42; scripts/bootstrap_layerwise_rho.py).

Trajectory inference. Diffusion pseudotime (scanpy DPT) was computed on PCA-reduced external embeddings (n_pcs = 30, n_neighbors = 30, n_dcs = 10), rooted at the LSC class, and correlated (Spearman) with the reference layer-ordering pseudotime for each arm (scripts/analyze_trajectory_continuity.py).

TAC sensitivity analysis. Because the TAC class is extremely rare (171/14,775 cells; 1.2%), donor-level macro-F1 was recomputed excluding TAC (three classes), with the same bootstrap procedure, to assess whether the rare class drives the primary endpoint (scripts/analyze_tac_sensitivity.py).

CNS-anchored marker-program validation. To test whether the measured axis carries published biology, per-cell marker-program scores were computed as the mean of z-scored log1p expression over literature-anchored gene programs: stemness (KRT15, KRT19, TP63), corneal differentiation (KRT12, KRT3), and cycling (PCNA, MCM2, MKI67, TOP2A) on GSE153515; basal (KRT5, KRT14, COL17A1) and differentiation (KRT1, KRT10) keratins and cycling (MKI67, PCNA) on both epidermal datasets (GSE147424 and GSE202352). Expression was read from the raw GSE153515 matrix (barcode-aligned to the 14,775 epithelial cells; the QC-filtered matrix drops MKI67) and from the processed epidermis objects; GSE202352 (Wiedemann et al. [41], Cell Reports 2023) was assembled from its 10x filtered feature-barcode matrices (11 of 12 samples; one sample's matrix is corrupt at the GEO source and was excluded, covering all four donors and all three body sites) with the same marker-axis quartile binning as GSE147424 (scripts/process_gse202352.py; scripts/run_gse202352_eval.py). Each program score was correlated (Spearman) with the arm's per-cell axis pseudotime (axis definition, above); uncertainty was quantified by nonparametric bootstrap (500 resamples; seed 42) and permutation nulls (500 label-permutation replicates) standardized as z-scores (scripts/bio_validation.py). Circularity disclosure: KRT15, KRT19, KRT12, KRT3, and MKI67 appear in the locked annotation marker sets, so class-level enrichment of the corresponding programs is partially annotation-derived; the axis measurement and the program correlation are nevertheless distinct quantities, and the independent markers NFIB, CDH3, PCNA, and MCM2 (absent from all annotation sets) corroborate the recovery on the cornea. On epidermis, the basal-keratin and cycling programs are absent from the marker-score definition used to construct the ordinal layer labels.

4.8. Axis definition revision (protocol deviation)

The axis-anchored and soft-Spearman arms (Arms 7–10) were initially trained against a marker-score differentiation axis. An evaluation of seed 42 exposed a training-axis versus evaluation-metric direction mismatch: the marker-score axis and the per-layer-centroid-PCA axis used by the evaluation metric (Secondary endpoints) aligned with cosine 0.470. On 2026-08-10 the axis definition was revised (protocol deviation, user-approved; changelog 2026-08-10-paper1-axis-refit-option-a) to a per-layer-centroid-PCA axis computed over all epithelial training cells, bit-identical to the evaluation metric's axis (corneal_fm/axis_ref.py, mode = "centroid_pca"; cache schema v2); the marker-score axis is retained as a reference mode. Arms 7–11 and the cross-tissue evaluation used the revised axis.

4.9. Cross-architecture frozen probes (scGPT-human, scFoundation, UCE)

To probe whether the ordinal-geometry measurements generalize across single-cell foundation models without incurring a second fine-tuning protocol, the frozen embeddings of scGPT-human (perturblab/scgpt-human checkpoint; 12-layer Transformer, 512-dimensional hidden state, continuous log1p input, CLS cell embedding, 51.86M parameters; scgpt 0.2.5, git commit cebd6fae) were evaluated on both external test sets. Cells were tokenized with the model's 60,697-gene vocabulary (append_cls = True, maximum sequence length 1,200

with the CLS token kept first) and embedded at the CLS position of the final layer; the differentiation-gradient ρ was computed with the same marker-gradient metric as all other arms, and class separability was measured with a four-class logistic-regression probe (cell-level 80/20 split, five repeats, mean macro-F1). Gene-symbol mapping for GSE153515 was verified against HGNC: 10,379 of 10,605 ENSEMBL identifiers resolved authoritatively, of which 32 of the 9,175 vocabulary-matched genes (0.35%) carry post-download HGNC renaming (none core corneal markers), with no effect on the input.

Frozen scFoundation (01B-resolution checkpoint; genbio-ai/scFoundation Hugging Face mirror, checkpoint sha256 9f40bf32…; xTrimoGene encoder–decoder, 119.25M parameters [59]) was embedded with the official inference pipeline (get_embedding.py; input_type singlecell, output_type cell, pool_type all, tgthighres a5, pre_normalized F, version ce), yielding per-cell embeddings of dimension 3,072 (concatenation of last-layer, second-to-last, max-pooled, and mean-pooled hidden states over the 19,264-gene vocabulary). Raw counts were aligned to the model's 19,264 HGNC-symbol vocabulary (9,162 of 9,175 input genes matched for GSE153515; 10,296 of 11,274 for GSE299152; unmatched genes zero-filled). scFoundation's pretraining corpus (over 50 million human cells) is not publicly enumerated, so corpus membership of the two test sets could not be verified and is reported as inferred out-of-corpus.

Frozen UCE (Universal Cell Embedding; 33-layer checkpoint 33l_8ep_1024t_1280.torch, Figshare 10.6084/m9.figshare.24320806, checkpoint sha256 3f458726…; [18]) cell embeddings were computed with the official inference script (eval_single_anndata.py; species human, sample_size 1,024 expression-weighted genes sampled per cell with replacement, pad_length 1,536, CLS-token hidden state of dimension 1,280, L2-normalized). Because UCE cell representations depend on stochastic gene sampling, each dataset was embedded three times (seeds 42, 123, 456) and the differentiation-gradient ρ and probe macro-F1 were averaged across repeats (per-repeat values in Table S8). UCE was pretrained on the CELLxGENE census 2023-04-25 release; GSE153515 is absent from the closest queryable census release (2023-05-15) and GSE299152 postdates the corpus, so both test sets are reported as out-of-corpus (inference).

4.10. CAFT promotion criteria

CAFT (Arm 4) was promoted over naive fine-tuning (Arm 3) only if all four criteria were satisfied simultaneously: (1) external continuity improvement (ρ or LAA) relative to Arm 3; (2) non-inferior classification (bootstrap 95% CI overlap or non-inferiority within a prespecified margin); (3) seed and donor consistency of the improvement direction; and (4) survival of the locked provenance audit and donor- and study-level disjointness verification. Failure of any single criterion retained Arm 3 as the preferred strategy, with CAFT reported as an unsuccessful candidate; no post hoc tuning or selective reporting was permitted.

4.11. Provenance and contamination audit

All computational artifacts were version-locked prior to analysis: model weights (retrieved 2026-07-23), training and test data files (downloaded 2026-07-22), and SHA-256 checksums for model weights and train/test AnnData were recorded, and the fair-split manifest was generated before any training. Because Geneformer was pretrained on Genecorpus-104M, we assessed whether corneal or ocular surface datasets from our training atlas might have been included

in the pretraining corpus. For Genecorpus-30M (V1; 561 datasets with a public manifest), all 43 eye-related datasets were identified; all were retinal or retinal pigment epithelium datasets, with zero corneal datasets, and all eight GEO accessions corresponding to our training studies returned no matches. For Genecorpus-104M (V2; 2,903 datasets), the official corpus source table (Supplementary table MOESM3, sheet `genecorpus_104m`; 563 unique DOIs) was downloaded from the Nature Computational Science supplementary materials and used for a DOI-level audit. GSE153515's primary publication DOI (10.1016/j.jtos.2020.12.004; Li et al., The Ocular Surface 2021) resolves to a row in the corpus table (53,290 cells, organ = eye); because all corpus datasets predate assembly and GSE153515 was public from 2021-01-14, no temporal exemption applies, and GSE153515 is reported as same-corpus (least conservative). GSE299152 [36] (GEO public 2025-11-18) postdates corpus assembly and its DOI is absent from the table, making it the post-cutoff, out-of-corpus (most conservative) test. The cross-tissue control GSE147424 [37] (He et al., J Allergy Clin Immunol 2020) likewise resolves to a row in the corpus table (organ = skin, 45,200 cells), so the epidermal test is within the pre-training corpus; its frozen differentiation axis is nevertheless weak ($\rho = 0.24$ versus 0.81 for the same-corpus corneal test), showing that corpus membership and axis encoding strength are decoupled (Results, three-domain boundary map). The second cross-tissue dataset, GSE202352 [41] (Wiedemann et al., Cell Reports 2023; organ = skin, 15,243 cells), is also within the corpus; both epidermal tests are therefore reported as in-corpus cross-tissue controls, and the cross-tissue reversal is interpreted as an axis-encoding effect rather than a corpus-absence effect. For the cross-architecture probe, scGPT-human was pretrained on the CELLxGENE census 2023-05-15 release (33,758,887 unique normal human cells; [13] Data Availability). GSE153515 is absent from both the CELLxGENE curation catalog and the census 2023-05-15 dataset manifest (verified via the CZ curation API and the census S3 h5ads directory), and GSE299152 (GEO released 2025-11-18) postdates the corpus; both external tests are therefore out-of-corpus for scGPT. A secondary finding: two training-atlas studies (lako_GSE155683, 9,972 cells; li_GSE147979, 423 cells) resolve to corpus DOIs (10,395 of 99,606 cells, 10.4%), disclosed for completeness; this does not affect the external-test conclusions, which rest on donor- and study-level disjointness from the training atlas. As a corroborating cell-level check in the spirit of scContam [60], bounded Signal A fingerprint sampling (0.54% of the corpus, 491,227 fingerprints; vectorized MinHash m = 128, b = 16, r = 8, banded LSH, exact Jaccard confirmation) found zero hits for the 15,827 GSE153515 cells (max Jaccard 0.0), with 100% positive-control recall and zero negative-control false positives; because full-corpus coverage (approximately 600 GB) was not attained, this run validates the pipeline but does not constitute exclusion evidence. The DOI-level audit is the primary evidence.

CAFT implementation audit. A post hoc audit established that the continuity term $\mathcal{L}_{\text{cont}}$ never executed: in Transformers 4.57 the `on_epoch_begin` hook is dispatched only to registered `TrainerCallback` objects, and `CAFTTrainer.on_epoch_begin` was a plain method that no callback ever invoked, so `layer_centroids` was never populated and every archived run trained pure cross-entropy (archived scripts byte-identical; training histories contain only the `loss` key). The $\lambda = 0.5$ checkpoints' higher final-layer ρ ($+0.210$ versus -0.008) is therefore run-to-run nondeterminism, not restoration (audit forensics and corrected ECE values in artifacts/paper1_pipeline_audit/). A correctly wired continuity objective was subsequently tested inside the ordinal-CAFT arm, where it executed genuinely but conferred

no measurable benefit beyond the ordinal head (Methods, Ordinal head formulation; Results, ordinal head).

4.12. Statistical analysis plan

All statistical analyses were prespecified in the locked evaluation configuration prior to model training. The primary comparison was between Arm 3 and Arm 4, with Arms 1 and 2 as reference baselines. All primary same-corpus claims were submitted to a single Benjamini–Hochberg FDR family of 26 comparisons (q threshold 0.05), spanning, per arm, gradient-positivity ($\rho > 0$), layer-adjacency-above-chance (LAA > chance), and classification-parity (F1 = naive F1) tests across the frozen, naive, CAFT, ordinal, AA-FT, soft-Spearman, and flow arms (scripts/fdr_nulls_analysis.py; artifacts/paper1_fdr_nulls/bh_fdr_table.json). Permutation nulls were constructed by label shuffling (500 replicates) and feature shuffling (300 replicates) of the differentiation-gradient measurement, and observed ρ values were standardized against the null mean and SD (z-scores; artifacts/paper1_fdr_nulls/null_controls.json). Uncertainty was quantified by nonparametric bootstrap resampling at the donor level (n_boot = 1,000; seed = 42). Model stability was assessed across three independent training seeds, reported as mean ± range. All hypothesis tests were two-sided; $p < 0.05$ was considered significant. No interim analyses or adaptive stopping rules were employed.

4.13. Software and reproducibility

All analyses were performed in Python 3.12, using Scanpy [61], harmonypy [54, 62], Hugging Face Transformers [63], PyTorch [64], and scikit-learn [65]. The CAFT training pipeline was implemented as a custom extension of the Hugging Face Trainer class (corneal_fm/caft.py); evaluation used the Tier-B module (corneal_fm/evaluation.py); external annotation used the locked pipeline scripts and configuration files. Gradient-preserving objective training used corneal_fm/axis_ref.py and scripts/method_arm_gf.py; the three-domain boundary evaluation used scripts/run_method_arms_eval.py, scripts/epidermis_tpoac_eval.py, and scripts/summarize_tpoac_boundary.py; FDR and permutation-null control used scripts/fdr_nulls_analysis.py; the CNS-anchored marker-program validation used scripts/bio_validation.py and scripts/extract_ordinal_embeddings.py.

Resource availability

Lead contact. Further information and requests for resources should be directed to and will be fulfilled by the lead contact, Qi Dai (dq@mail.eye.ac.cn).

Materials availability. This study did not generate new unique reagents or biological materials; all data are public (below).

Data and code availability.

- **Training data.** CELLxGENE Discover collection 0f7d022a-46c7-4e64-be4c-e34adbb78089 (Single cell atlas of the human ocular surface; Human Ocular Surface Atlas), current dataset-version of dataset 8ef06ea4-8d0f-4303-a652-7d4604402b1f, distributed as file 30127f02-d0b9-4b56-bf4c-e5b77e4989bd.h5ad; MD5 39f098367af95bc3e1ee92528ceb251e.

- **External test data.** GEO GSE153515 [3]; 17,587 cells from two donors (GSM4646295, GSM4646296). GEO GSE299152 [36]; 21,374 epithelial cells from three donors after QC. Cross-tissue: GEO GSE147424 [37].

- **Cross-architecture probes.** scGPT-human checkpoint (perturblab/scgpt-human; scgpt 0.2.5, git commit cebd6fae); scFoundation 01B-resolution checkpoint (genbio-ai/scFoundation Hugging Face mirror); UCE 33-layer checkpoint (Figshare 10.6084/m9.figshare.24320806); frozen-probe diagnostics in artifacts/scgpt_repro/frozen_diag/, artifacts/scfoundation_repro/frozen_diag/, and artifacts/uce_repro/frozen_diag/.

- **Model.** Geneformer-V2-104M, Hugging Face model hub (ctheodoris/Geneformer; snapshot commit 04c2b2e84da7c0f385c3f9ad8f3ec24bab6650e5).

- **Code.** Analysis code is available at `https://github.com/Xu-Kunk/Cornea-Ai` and will be made publicly accessible upon acceptance.

- **Evaluation scripts.** scripts/run_tier_b_eval.py; scripts/run_tier_b_baselines.py; scripts/measure_layerwise_rho.py; scripts/bootstrap_layerwise_rho.py; scripts/analyze_trajectory_continuity.py; scripts/analyze_tac_sensitivity.py; scripts/protocol_15b_annotate_gse299152.py; scripts/method_arm_gf.py; scripts/run_method_arms_eval.py; scripts/epidermis_tpoac_eval.py; scripts/process_gse202352.py; scripts/run_gse202352_eval.py; scripts/summarize_tpoac_boundary.py; scripts/fdr_nulls_analysis.py; scripts/bio_validation.py; scripts/extract_ordinal_embeddings.py.

- **Results.** artifacts/paper1_tier_b_eval/*.json (six fine-tuned-arm and two baseline-arm result files); artifacts/paper1_tier_b_eval_gse299152/*.json (out-of-corpus results); artifacts/paper1_ordinal_eval/*.json (ordinal arms); artifacts/paper1_epidermis_eval/epidermis_rho_results.json (cross-tissue, GSE147424); artifacts/paper1_epidermis_eval_gse202352/gse202352_rho_results.json (cross-tissue, GSE202352); artifacts/paper1_bio_validation/bio_validation_results.json (CNS-anchored marker-program validation); artifacts/paper1_fdr_nulls/ (bh_fdr_table.json, null_controls.json, contamination_audit_protocol.md).

- **Reproducibility.** All experiments are reproducible with fixed random seeds (42, 123, 456), locked configuration files (config/external_test_gse153515.yml, config/external_test_gse299152.yml), and the SHA-256-locked fair-split manifest.

References

[1] G. Cotsarelis, S.-Z. Cheng, G. Dong, T.-T. Sun, R. M. Lavker, Existence of slow-cycling limbal epithelial basal cells that can be preferentially stimulated to proliferate: implications on epithelial stem cells, Cell 57 (2) (1989) 201–209. doi:10.1016/0092-8674(89)90958-6.

[2] H. S. Dua, A. Azuara-Blanco, Limbal stem cells of the corneal epithelium, Survey of Ophthalmology 44 (5) (2000) 415–425. doi:10.1016/S0039-6257(00)00109-0.

[3] D.-Q. Li, S. Kim, J.-M. Li, Q. Gao, J. Choi, F. Bian, J. Hu, Y. Zhang, R. Lu, Y. Lu, S. C. Pflugfelder, Single-cell transcriptomics identifies limbal stem cell population and cell types mapping its differentiation trajectory in limbal basal epithelium of human cornea, The Ocular Surface 20 (2021) 20–32. doi:10.1016/j.jtos.2020.12.004.

[4] J. Collin, R. Queen, D. Zerti, B. Dorgau, M. Georgiou, I. Djidrovski, R. Hussain, J. M. Coxhead, A. Joseph, P. Rooney, S. Lisgo, F. Figueiredo, L. Armstrong, M. Lako, A single cell atlas of human cornea that defines its development, limbal progenitor cells and their interactions with the immune cells, The Ocular Surface 21 (2021) 279–298. doi:10.1016/j.jtos.2021.03.010.

[5] P. Català, N. Groen, J. A. Dehnen, J. S. Soares, A. J. H. van Velthoven, R. M. M. A. Nuijts, M. M. Dickman, V. L. S. LaPointe, Single cell transcriptomics reveals the heterogeneity of the human cornea to identify novel markers of the limbus and stroma, Scientific Reports 11 (1) (2021) 21727. doi:10.1038/s41598-021-01015-w.

[6] J.-M. Li, S. Kim, D.-Q. Li, Q. Gao, J. Choi, F. Bian, J. Hu, Y. Zhang, R. Lu, S. C. Pflugfelder, Single-cell transcriptomics identifies a unique entity and signature markers of transit-amplifying cells in human corneal limbus, Investigative Ophthalmology & Visual Science 62 (9) (2021) 36. doi:10.1167/iovs.62.9.36.

[7] A. J. Ligocki, et al., Molecular characteristics and spatial distribution of adult human corneal cell subtypes, Scientific Reports 11 (2021) 16323. doi:10.1038/s41598-021-94933-8.

[8] G. Maiti, M. Alameen, D. Nayak, M. S. Sridhar, S. Ghosh, Single cell rna-seq of human cornea organoids identifies cell fates of a developing immature cornea, PNAS Nexus 1 (5) (2022) pgac246. doi:10.1093/pnasnexus/pgac246.

[9] R. Chen, CZI Cell Science / Human Cell Atlas Eye Network, Single-cell atlas of the human ocular surface, CELLxGENE Discover collection 0f7d022a-46c7-4e64-be4c-e34adbb78089, accessed 2026. Pre-publication consortium atlas; cited as a collection (2026).
URL https://cellxgene.cziscience.com/collections/0f7d022a-46c7-4e64-be4c-e34adbb78089

[10] J. A. Arts, S. Fallo, M. S. Florencio, J. G. A. Smits, D. Lima Cunha, J. A. Y. Roubroeks, et al., Prediction of cell states and key transcription factors of the human cornea through integrated single-cell omics analyses, PNAS Nexus 4 (8) (2025) pgaf235. doi:10.1093/pnasnexus/pgaf235.

[11] C. V. Theodoris, L. Xiao, A. Chopra, M. D. Chaffin, Z. R. Al Sayed, M. C. Hill, H. Mantineo, E. M. Brydon, Z. Zeng, X. S. Liu, P. T. Ellinor, Transfer learning enables predictions in network biology, Nature 618 (7965) (2023) 616–624. doi:10.1038/s41586-023-06139-9.

[12] H. Chen, et al., Scaling and quantization of large-scale foundation model enables resource-efficient predictions in network biology, Nature Computational Science 6 (5) (2026) 450–463. doi:10.1038/s43588-026-00972-4.

[13] H. Cui, C. Wang, H. Maan, K. Pang, F. Luo, N. Duan, B. Wang, scgpt: toward building a foundation model for single-cell multi-omics using generative ai, Nature Methods 21 (8) (2024) 1470–1480. doi:10.1038/s41592-024-02201-0.

[14] F. Yang, W. Wang, F. Wang, Y. Fang, D. Tang, J. Huang, H. Lu, J. Yao, scbert as a large-scale pretrained deep language model for cell type annotation of single-cell rna-seq data, Nature Machine Intelligence 4 (10) (2022) 852–866. doi:10.1038/s42256-022-00534-z.

[15] S. Hou, P. Yang, W. Ma, J.-X. Wang, X. Zhou, A unified framework enables accessible deployment and comprehensive benchmarking of single-cell foundation models, bioRxiv preprint (2026). doi:10.64898/2026.01.06.698060.

[16] J. Wu, Q. Ye, Y. Wang, et al., Biology-driven insights into the power of single-cell foundation models, Genome Biology 26 (2025) 334. doi:10.1186/s13059-025-03781-6.

[17] T. Liu, K. Li, Y. Wang, H. Li, H. Zhao, Evaluating the utilities of foundation models in single-cell data analysis, Advanced Science 13 (27) (2026) e14490. doi:10.1002/advs.202514490.

[18] Y. Rosen, Y. Roohani, A. Agrawal, L. Samotorčan, Tabula Sapiens Consortium, S. R. Quake, J. Leskovec, Universal cell embedding provides a foundation model for cell biology, Nature 656 (8126) (2026) 183–191. doi:10.1038/s41586-026-10689-z.

[19] P. Qiu, Q. Chen, H. Qin, S. Fang, Y. Zhang, Y. Zhang, et al., Biollm: A standardized framework for integrating and benchmarking single-cell foundation models, Patterns 6 (8) (2025) 101326. doi:10.1016/j.patter.2025.101326.

[20] T. Liu, T. Chen, W. Zheng, X. Luo, Y. Chen, H. Zhao, Embeddings from language models are good learners for single-cell data analysis, Patterns 7 (2) (2026) 101431. doi:10.1016/j.patter.2025.101431.

[21] C. Wang, Y. T. Chen, et al., Evaluating large language models for evidence-based clinical question answering, Patterns 7 (5) (2026) 101519. doi:10.1016/j.patter.2026.101519.

[22] Z. Ma, J. Li, J. Hu, B. Yu, Uncovering smooth structures in single-cell data with neighbor embeddings guided by predictability-computability-stability, Patterns 7 (7) (2026) 101621. doi:10.1016/j.patter.2026.101621.

[23] F. He, R. Fei, J. E. Krull, Y. Yu, X. Zhang, X. Wang, et al., Harnessing the power of single-cell large language models with parameter-efficient fine-tuning using scpeft, Nature Machine Intelligence 8 (2026) 118–133. doi:10.1038/s42256-025-01170-z.

[24] S. Weerasekara, N. Darras, S. Kamarthi, C. Price, J. Isaacs, Prototype guided post-pretraining for single-cell representation learning, arXiv preprint arXiv:2605.07938 (2026).

[25] X. Zhou, Z. Wang, Y. Ling, Q. Tian, Z. Zhang, Y. Li, L. Chen, P. Zhou, Benchmarking zero-shot single-cell foundation model embeddings for cellular dynamics reconstruction, bioRxiv preprint (2026). doi:10.64898/2026.03.10.710748.

[26] S. Pan, E. Withnell, C. Celik, M. Secrier, From cell types to cell states: adapting foundation models for cellular plasticity, bioRxiv preprint, v2 January 2026 (2024). doi:10.1101/2024.08.16.608311.

[27] I. Kendiukhov, et al., Attention layers in single-cell foundation models organize as protein–protein interactions in early layers and transcriptional regulation in late layers, BMC Genomics 27 (2026) 12965. doi:10.1186/s12864-026-12965-8.

[28] I. Kendiukhov, et al., Residual-stream geometry in single-cell foundation models, BMC Bioinformatics 27 (2026) 6538. doi:10.1186/s12859-026-06538-5.

[29] S. Ding, J. Li, R. Luo, H. Cui, B. Wang, R. Chen, scgpt: end-to-end protocol for fine-tuned retinal cell type annotation, Nature Protocols 21 (3) (2026) 873–893. doi:10.1038/s41596-025-01220-1.

[30] M. Teles, et al., Unraveling learning characteristics of transformer models for molecular design, Patterns 6 (12) (2025) 101392. doi:10.1016/j.patter.2025.101392.

[31] W. A. Van Amsterdam, et al., When accurate prediction models yield harmful self-fulfilling prophecies, Patterns 6 (4) (2025) 101229. doi:10.1016/j.patter.2025.101229.

[32] J. Burridge, et al., The inadequacy of offline large language model evaluations, Patterns 6 (12) (2025) 101397. doi:10.1016/j.patter.2025.101397.

[33] C. Ma, T. Obuchi, T. Tanaka, Neural collapse in cumulative link models for ordinal regression: An analysis with unconstrained feature model, arXiv preprint, neurIPS 2025 (camera-ready); arXiv:2506.05801 (2025).

[34] V. Y. Civale, R. Semeraro, A. D. Bagdanov, A. Magi, Intermediate layers encode optimal biological representations in single-cell foundation models, arXiv preprint, arXiv:2604.14838 (2026).

[35] K. Inecik, A. Rose, M. Haniffa, M. Luecken, F. J. Theis, Beyond visual inspection: Principled benchmarking of single-cell trajectory representations with sctram, bioRxiv preprint (2025). doi:10.1101/2025.06.23.661141.

[36] R. Martínez-Carrasco, P. Argüeso, Single-cell mapping of the glycocalyx in the adult human and mouse ocular surface, PNAS Nexus 4 (11) (2025) pgaf348. doi:10.1093/pnasnexus/pgaf348.

[37] H. He, H. Suryawanshi, P. Morozov, J. Gay-Mimbrera, E. Del Duca, H. J. Kim, N. Kameyama, Y. Estrada, E. Der, J. G. Krueger, J. Ruano, T. Tuschl, E. Guttman-Yassky, Single-cell transcriptome analysis of human skin identifies novel fibroblast subpopulation and enrichment of immune subsets in atopic dermatitis, Journal of Allergy and Clinical Immunology 145 (6) (2020) 1615–1628. doi:10.1016/j.jaci.2020.01.042.

[38] M. Farazi, S. Ramasinghe, B. S. Ciftler, M. A. Turza, S. Rahman, The gate always closes: On injecting auxiliary signals into frozen vision-language models, arXiv preprint, arXiv:2607.23335 (2026).

[39] X. Shi, W. Cao, S. Raschka, Deep neural networks for rank-consistent ordinal regression based on conditional probabilities, Pattern Analysis and Applications (2023). doi:10.1007/s10044-023-01181-9.

[40] E. Fuchs, H. Green, Changes in keratin gene expression during terminal differentiation of the keratinocyte, Cell 19 (4) (1980) 1033–1042. doi:10.1016/0092-8674(80)90094-X.

[41] J. Wiedemann, A. C. Billi, F. Bocci, G. Kashgari, E. Xing, L. C. Tsoi, L. Meller, W. R. Swindell, R. Wasikowski, X. Xing, F. Ma, M. Gharaee-Kermani, J. M. Kahlenberg, P. W. Harms, E. Maverakis, Q. Nie, J. E. Gudjonsson, B. Andersen, Differential cell composition and split epidermal differentiation in human palm, sole, and hip skin, Cell Reports 42 (1) (2023) 111994. doi:10.1016/j.celrep.2023.111994.

[42] V. Papyan, X. Y. Han, D. L. Donoho, Prevalence of neural collapse during the terminal phase of deep learning training, Proceedings of the National Academy of Sciences 117 (40) (2020) 24652–24663. doi:10.1073/pnas.2015509117.

[43] A. Kumar, A. Raghunathan, R. Jones, T. Ma, P. Liang, Fine-tuning can distort pre-trained features and underperform out-of-distribution, arXiv preprint arXiv:2202.10054, published at ICLR 2022 (2022).

[44] G. Andriopoulos, Z. Dong, L. Guo, Z. Zhao, K. Ross, The prevalence of neural collapse in neural multivariate regression, in: Advances in Neural Information Processing Systems 37 (NeurIPS 2024), Vol. 37, 2024, pp. 126417–126451.

[45] A. Aghajanyan, A. Shrivastava, A. Gupta, N. Goyal, L. Zettlemoyer, S. Gupta, Better fine-tuning by reducing representational collapse, arXiv preprint, iCLR 2021; arXiv:2008.03156 (2021).

[46] K. Z. Kedzierska, L. Crawford, A. P. Amini, A. X. Lu, Zero-shot evaluation reveals limitations of single-cell foundation models, Genome Biology 26 (2025) 101. doi:10.1186/s13059-025-03574-x.

[47] Y. Zeng, et al., Cellfm: a large-scale foundation model pre-trained on transcriptomics of 100 million human cells, Nature Communications 16 (2025) 4679. doi:10.1038/s41467-025-59926-5.

[48] S. Cultrera di Montesano, et al., Improving atlas-scale single-cell annotation models with hierarchical cross-entropy loss, Nature Computational Science 6 (3) (2026) 243–249. doi:10.1038/s43588-025-00945-z.

[49] C. Hu, C. Y.-C. Chen, Artificial intelligence virtual cells: From measurements to decisions across modality, scale, dynamics, and evaluation, arXiv preprint, arXiv:2510.12498 (2025). doi:10.48550/arXiv.2510.12498.

[50] C. Bunne, Y. Roohani, Y. Rosen, A. Gupta, X. Zhang, M. Roed, T. Alexandrov, M. AlQuraishi, P. Brennan, D. B. Burkhardt, et al., How to build the virtual cell with artificial intelligence: Priorities and opportunities, Cell (2024). doi:10.1016/j.cell.2024.11.015.

[51] L. Weidener, M. Brkić, M. Jovanović, E. Ulgac, A. Meduri, Vcbench: A multi-dimensional benchmark for single-cell foundation models, bioRxiv preprint (2026). doi:10.64898/2026.06.18.733146.

[52] S. Dou, Q. Wang, X. Qi, B. Zhang, H. Jiang, S. Chen, J. Li, M. Li, Molecular identity of human limbal heterogeneity involved in corneal homeostasis and privilege, The Ocular Surface 21 (2021) 206–220. doi:10.1016/j.jtos.2021.04.010.

[53] P. Gautam, K. Hamashima, Y. Chen, Y. Zeng, B. Makovoz, B. H. Parikh, Y. Lee, K. A. Lau, B. Lim, B. X. Ng, et al., Multi-species single-cell transcriptomic analysis of ocular compartment regulons, Nature Communications 12 (1) (2021) 5675. doi:10.1038/s41467-021-25968-8.

[54] I. Korsunsky, N. Millard, J. Fan, K. Slowikowski, F. Zhang, K. Wei, Y. Baglaenko, M. Brenner, P.-R. Loh, S. Raychaudhuri, Fast, sensitive and accurate integration of single-cell data with harmony, Nature Methods 16 (12) (2019) 1289–1296. doi:10.1038/s41592-019-0619-0.

[55] Y. Du, F. Fan, S. Zhao, J. Cao, Q. Lin, K. He, T. Liu, B. Qin, M. Feng, Anchoring refusal direction: Mitigating safety risks in tuning via projection constraint, arXiv preprint, arXiv:2509.06795 (2025).

[56] tadteo, Abag-rank: attribute-based augmentation for differentiable ranking (soft-spearman implementation), GitHub repository, accessed 2026 (2025).
URL https://github.com/tadteo/ABAG-Rank

[57] E. Ramzi, N. Audebert, C. Rambour, A. Araujo, X. Bitot, N. Thome, Optimization of rank losses for image retrieval, IEEE Transactions on Pattern Analysis and Machine IntelligenceArXiv:2309.08250 (2025).

[58] M. Blondel, O. Teboul, Q. Berthet, J. Djolonga, Fast differentiable sorting and ranking, arXiv preprint arXiv:2002.08871 (2020).

[59] M. Hao, J. Gong, X. Zeng, C. Liu, Y. Guo, X. Cheng, T. Wang, J. Ma, X. Zhang, L. Song, Large-scale foundation model on single-cell transcriptomics, Nature Methods (2024). doi:10.1038/s41592-024-02305-7.

[60] S. Ali, Auditing pretraining contamination in single-cell foundation model benchmarks, arXiv preprint, arXiv:2607.20572 (2026).

[61] F. A. Wolf, P. Angerer, F. J. Theis, Scanpy: large-scale single-cell gene expression data analysis, Genome Biology 19 (1) (2018) 15. doi:10.1186/s13059-017-1382-0.

[62] K. Slowikowski, Korsunsky, Ilya, harmonypy: Python implementation of harmony, GitHub repository, accessed 2026 (2020).
URL https://github.com/slowkow/harmonypy

[63] T. Wolf, L. Debut, V. Sanh, J. Chaumond, C. Delangue, A. Moi, P. Cistac, T. Rault, R. Louf, M. Funtowicz, J. Davison, S. Shleifer, P. von Platen, C. Ma, Y. Jernite, J. Plu, C. Xu, T. Le Scao, S. Gugger, M. Drame, Q. Lhoest, A. M. Rush, Huggingface's transformers: State-of-the-art natural language processing, in: Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: System Demonstrations, 2020, pp. 38–45. doi:10.18653/v1/2020.emnlp-demos.6.

[64] A. Paszke, S. Gross, F. Massa, A. Lerer, J. Bradbury, G. Chanan, T. Killeen, Z. Lin, N. Gimelshein, L. Antiga, A. Desmaison, A. Kopf, E. Yang, Z. DeVito, M. Raison, A. Tejani, S. Chilamkurthy, B. Steiner, L. Fang, J. Bai, S. Chintala, Pytorch: An imperative style, high-performance deep learning library, in: Advances in Neural Information Processing Systems 32, 2019, pp. 8026–8037.

[65] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, É. Duchesnay, Scikit-learn: Machine learning in python, Journal of Machine Learning Research 12 (2011) 2825–2830.
URL http://jmlr.org/papers/v12/pedregosa11a.html