Fine-tuning collapses differentiation gradients in single-cell foundation models in a domain-specific manner

Kunhui Xu[1], Chenchen Wang[1], and Qi Dai[1,*]

[1]National Clinical Research Center for Ocular Diseases, Eye Hospital, Wenzhou Medical University, Wenzhou, Zhejiang 325035, PR China
[*]Correspondence: dq@mail.eye.ac.cn

SUMMARY

Supervised fine-tuning of single-cell foundation models is routine, yet whether it preserves the continuous differentiation geometry of pretrained representations has not been quantified. Using the human corneal epithelium, a four-stage differentiation continuum, we show that classification fine-tuning collapses the differentiation gradient in a cross-architecture spectrum. Collapse is corpus-determined: on same-corpus tissue, Geneformer degrades the gradient (Spearman ρ 0.806 to 0.096), and scGPT-human shows moderate loss (0.789 to 0.44, best-val), whereas out-of-corpus and cross-tissue effects remain mild or reversed. Collapse determinism is architecture-determined: Geneformer collapses deterministically, while scGPT-human enters a seed-dependent bistable state. An ordinal (CORN) head rescues the gradient in both architectures; combined with LoRA, CORN+LoRA achieves a three-seed mean ρ of 0.924. We release a contamination-aware benchmark pairing same-corpus and out-of-corpus tests.

KEYWORDS

single-cell foundation models, classification fine-tuning, differentiation gradient, ordinal regression, contamination-aware benchmark, corneal epithelium, parameter-efficient adaptation, cross-architecture evaluation

INTRODUCTION

The corneal epithelium is maintained by a tightly regulated differentiation continuum from quiescent limbal stem cells (LSCs) at the limbal palisades of Vogt, through transit-amplifying cells (TACs), to committed basal epithelial cells, and ultimately to terminally differentiated squamous cells forming the protective ocular surface barrier[1]. This ordered progression is functionally required: disruption at any stage (LSC deficiency, aberrant amplification, or premature terminal differentiation) compromises corneal transparency and may lead to irreversible vision loss[2].

Single-cell RNA sequencing (scRNA-seq) has resolved this continuum at single-cell resolution, with marker genes for LSCs (e.g., *TCF4*, *POSTN*, *ABCG2*, and *TSPAN7*[3]; *TP63* and *KRT15*[4,5]), TACs (e.g., *UBE2C*, *CDC20*, and *CENPF*[6]; *MKI67* and *BIRC5*[3]), and differentiated populations (e.g., *KRT12* and *KRT3*[4,7,8]) identified across independent studies and aggregated in community atlases such as the human ocular surface atlas in CELLxGENE Discover[9]. Yet these efforts produce study-specific embeddings that are not directly comparable across datasets, and no unified reference model maps individual cells to a shared, ordered differentiation-state space. The closest approach, a corneal cell-state reference model with support-vector-machine classification (cPredictor)[10], provides a supervised classifier rather than an evaluation of foundation-model representations or their geometric structure.

Foundation models pretrained on massive single-cell corpora offer a potential solution. Geneformer, a 12-layer Transformer pretrained on approximately 104 million cells from 2,903 datasets[11,12], and related architectures such as scGPT[13] and scBERT[14], learn contextual gene–gene representations that transfer across tissues. Systematic evaluations have tempered expectations: a unified benchmark of thirteen single-cell foundation models found classical PCA remains competitive in several settings, with pretrained embeddings advantageous mainly in low-label and transfer-learning scenarios[15]; biology-driven re-analyses questioned the pre-train-then-fine-tune paradigm[16]; and large-scale evaluations found that, for cell-type annotation, frozen or zero-shot usage generally underperforms fine-tuning, and freezing pretrained layers was not recommended[17]. Conversely, foundation models designed for zero-shot use can organize developmental lineages and differentiation hierarchies without task-specific training[18], indicating that pretrained representations can encode continuous differentiation structure. Benchmark-driven evaluation has become an established genre in this journal, spanning single-cell foundation models[19,20] and clinical question answering[21], and representation-quality assessment is emerging as a distinct concern in embedding evaluation[22]. Whether general-purpose representations adequately capture the ordered structure of a specific differentiation hierarchy, such as the corneal continuum, remains open; conversely, task-specific fine-tuning may improve discriminative accuracy while inadvertently collapsing the continuous geometric relationships among adjacent differentiation states that are biologically essential.

The fine-tuning literature has focused on whether adaptation improves classification accuracy and whether it causes forgetting of learned cell-type knowledge; full fine-tuning induces catastrophic forgetting at the label level, motivating parameter-efficient remedies[23]; a complementary line reports bounded fine-tuning gains and proposes marker-gene-guided post-pretraining instead[24]. With respect to geometry, a zero-shot benchmark reported geometric collapse of foundation-model embeddings during cellular dynamics reconstruction, with HVG/PCA outperforming them for trajectory recovery[25], and EMP-FM showed that naive fine-tuning often fails to resolve intermediate cell states under a dual classification–pseudotime objective[26]. Layer-resolved dissections of single-cell foundation models have begun to probe where biological information resides within the network[27,28], and fine-tuned representations are already deployed in production annotation protocols without quantifying whether adaptation preserves ordinal structure[29]; a related cautionary literature documents cases in which well-performing models fail in settings that matter[30,31]. What remains absent is a quantified measurement, under controlled fine-tuning conditions, of whether supervised adaptation preserves or destroys the continuous differentiation-state gradient that orders a hierarchy, which we term the *continuity problem*. The corneal epithelium is an ideal test system: its four ordered stages have well-established markers, so the expected ordinal structure is known *a priori* and gradient preservation can be measured directly in embedding space.

Here we quantify the effect of classification fine-tuning on differentiation-state geometry in the normal human corneal epithelium. Using a donor-disjoint fair-split protocol on a 10-study corneal and ocular surface atlas (99,606 training cells from 49 donors; 20,394 test cells from 25 donors; zero donor overlap), we compare two reference arms (HVG/PCA with a linear classifier; frozen Geneformer-V2-104M with a linear probe) with adapted configurations sharing identical preprocessing, data splits, classification head, and training budget: naive cross-entropy fine-tuning, a continuity-regularized objective (CAFT), an ordinal head (CORN, with and without the continuity term), and three further gradient-preserving objectives (AA-FT, soft-Spearman, flow), evaluated on a same-corpus donor-disjoint holdout (GSE153515; 17,587 cells, 2 donors), a truly out-of-corpus external test (GSE299152; 21,374 epithelial cells, 3 donors), and a cross-tissue epidermal test set. Our contributions are: (1) a quantified, three-domain measurement of fine-tuning-induced differentiation-gradient collapse: catastrophic on same-corpus data, mild on truly out-of-corpus data, and reversed on cross-tissue data; (2) an ordinal head reparameterization (CORN) as the

recipe preserving the gradient at matched classification accuracy, and only within the pretraining corpus; (3) a systematic evaluation of five alternative gradient-preserving objectives (three honest negatives, one partial, one unstable), in the spirit of recent critical-evaluation studies in this journal[31,32]; (4) a contamination-aware scFM benchmark for the normal human corneal epithelium, pairing a same-corpus donor-disjoint holdout with a truly out-of-corpus external test under a locked multi-seed protocol with BH-FDR control and permutation nulls; and (5) cross-architecture generalization of the collapse and the CORN rescue to scGPT-human and CellPLM, and the demonstration that CORN and LoRA are complementary: ordinal prior plus restricted overwrite yields the best-performing geometry-preserving recipe.

RESULTS

Fine-tuning collapses the differentiation gradient while preserving layer topology

Classification fine-tuning collapsed the differentiation gradient of the corneal epithelium while leaving layer topology intact. Two non-adapted references (HVG/PCA with a linear classifier; frozen Geneformer-V2-104M with a linear probe) were compared against end-to-end cross-entropy fine-tuning and a continuity-regularized variant (CAFT), all sharing identical preprocessing, data splits, classification head, and training budget, on a same-corpus, donor-disjoint holdout (GSE153515; Figure 1; three independent seeds per adapted arm; per-seed values in Table S1).

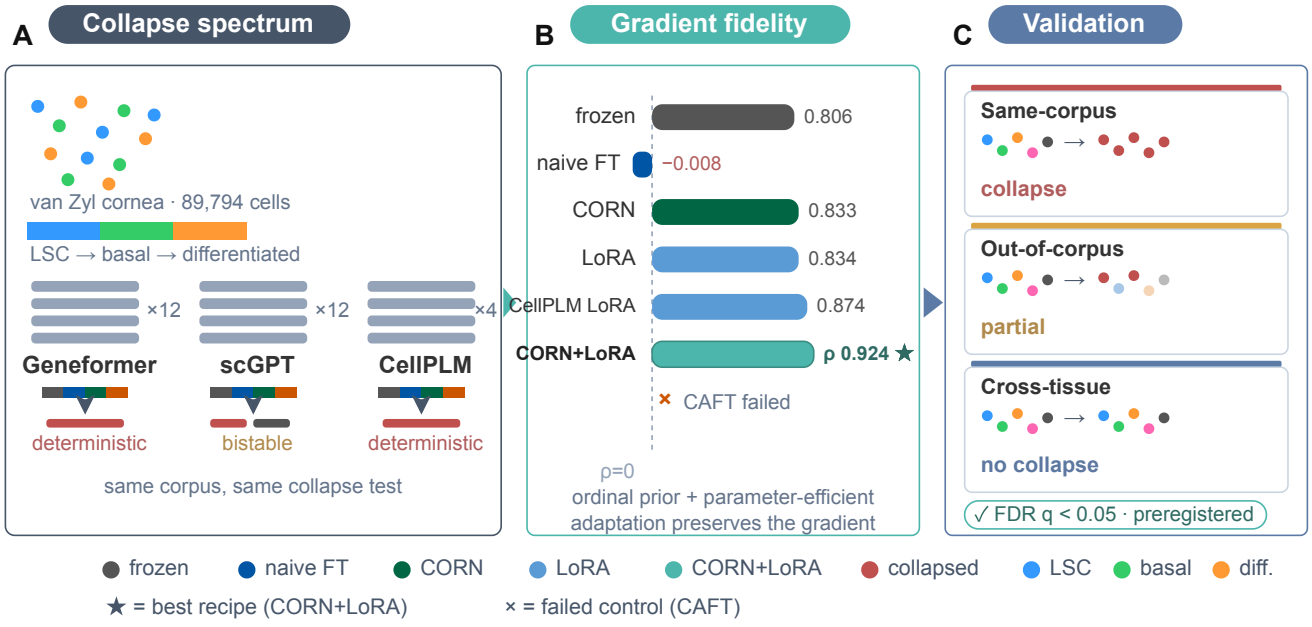


Figure 1: Study pipeline. Donor-disjoint training atlas (CELLxGENE), locked evaluation protocol, softmax-head and ordinal-head fine-tuning arms, and paired same-corpus (GSE153515) and post-cutoff out-of-corpus (GSE299152) external tests.

The two non-adapted reference arms preserved strong differentiation-gradient rank correlation on the same-corpus holdout: Spearman ρ was 0.821 for HVG/PCA and 0.806 for frozen Geneformer (Table 1). By contrast, naive supervised fine-tuning collapsed this ordinal signal to near zero: per-seed ρ ranged from -0.042 to $+0.054$ across seeds (three-seed mean -0.008), with bootstrap 95% confidence intervals spanning zero in every run (Figure 2, Table 1). The sign of ρ

was inconsistent across seeds: the residual ordinal correlation after fine-tuning was indistinguishable from noise. A randomized-head ablation – replacing the trained classifier with an untrained head before re-embedding – leaves the hidden states unchanged, ruling out that the collapse is a readout artifact of the trained classifier (Methods); the erasure resides in the fine-tuned body weights. The biological consequence of this collapse is quantified below (Figure 3).

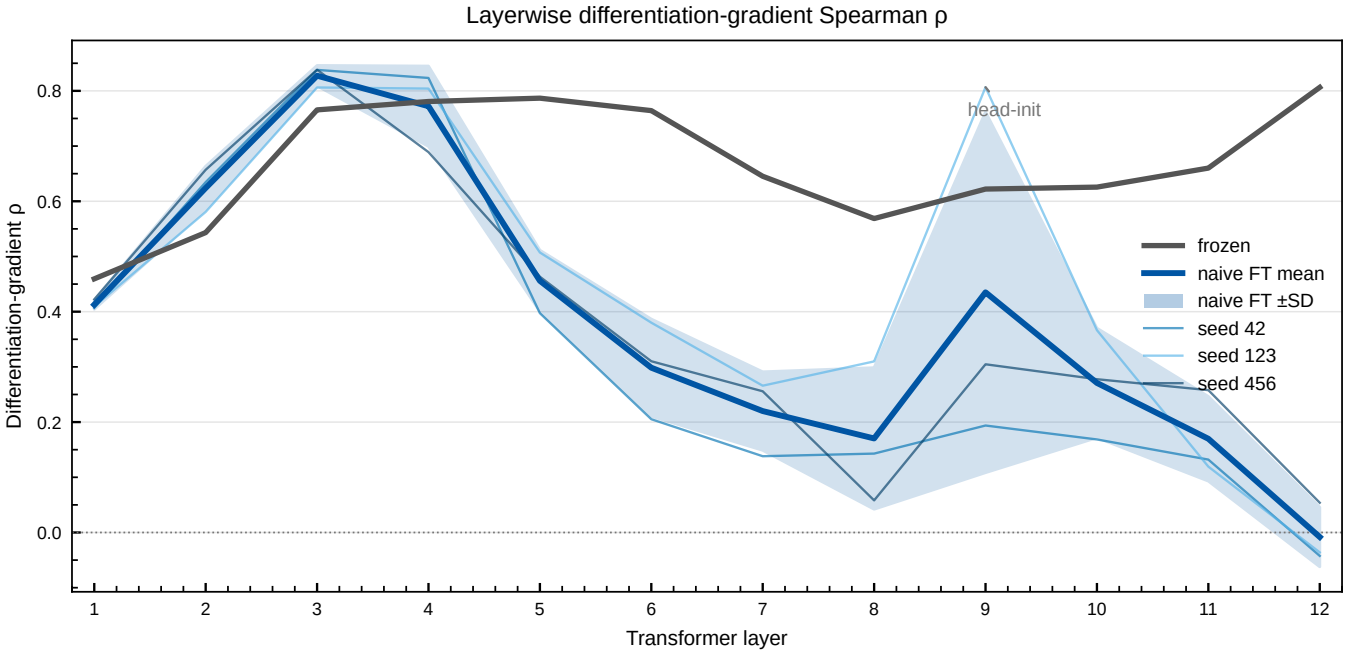
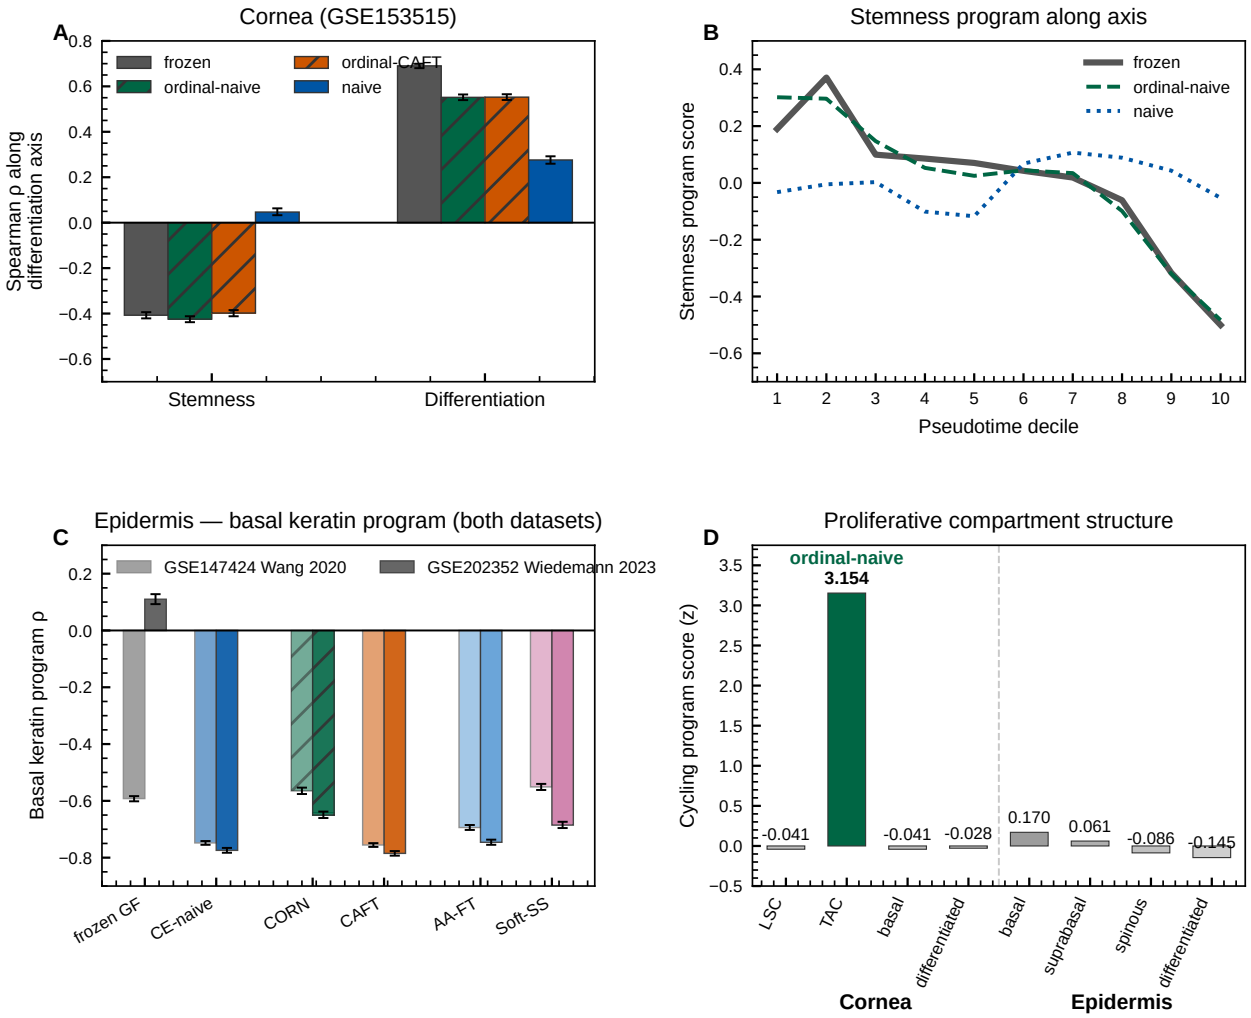


Figure 2: Layer-resolved analysis localizes the collapse to upper transformer layers. Layer-resolved differentiation-gradient Spearman ρ across hidden states L1–L12 for the Geneformer-based configurations (n = 14,775 external-test epithelial cells; three seeds per adapted arm; L0 excluded: constant [CLS]). Color key: thick dark gray line, frozen Geneformer; blue family, naive fine-tuning (seeds 42/123/456). The collapse is head-proximal: early layers retain the gradient, while L5–L12 decay. See also Figure S7 and Table S5.

Despite this gradient collapse, discrete layer topology was preserved. Layer adjacency accuracy (LAA) remained high in all four arms: 0.977 (HVG/PCA), 0.967 (frozen Geneformer), 0.977 (naive, three-seed mean), and 0.977 (CAFT, three-seed mean). Classification fine-tuning therefore disrupted the continuous ordering of the differentiation continuum while leaving local layer-adjacency structure essentially intact: continuous gradient geometry and discrete layer topology are dissociable under task adaptation.

The continuity-regularized arm must be read as a failed control. A post hoc audit of the archived training code established that the continuity term never executed in any run, because of an implementation defect in the training hook, so those checkpoints trained pure cross-entropy, identical in objective to the naive arm (Methods, Provenance and contamination audit). The $+0.210$ separation in final-layer ρ therefore reflects run-to-run nondeterminism rather than a causal effect of the regularization, and classification performance was likewise indistinguishable: macro-F1 was 0.645 (95% CI 0.616–0.674) for the continuity-regularized group versus 0.649 (95% CI 0.623–0.675) for naive fine-tuning ($\Delta = -0.004$; overlapping intervals). Throughout, we report this arm as an honest negative implementation result rather than a promoted safeguard. Fine-tuning nonetheless achieved the highest classification accuracy while collapsing the gradient to noise: differentiation geometry and classification performance are separable.

Layer-resolved analysis localized the collapse to upper transformer layers. To determine whether fine-tuning erases the differentiation-gradient signal globally or damages it in a layer-



Panels B and D: single run (no error bars).

Figure 3: The differentiation axis recovers CNS-published biology. (A) Cornea (GSE153515, n = 14,775): Spearman ρ of stemness and differentiation programs along each arm's axis (frozen with bootstrap 95% CI; adapted arms three-seed mean $\pm$ SD). (B) Stemness-program score by pseudotime decile: the graded decline is preserved by frozen and ordinal embeddings and flattened by naive fine-tuning. (C) Epidermis, both datasets (GSE147424, n = 18,355; GSE202352, n = 14,461): basal-keratin program ρ by arm; every adapted arm recovers the basal-off switch, reproducing the cross-tissue reversal. (D) Cycling program by class: the proliferative transient-amplifying compartment (corneal TAC; epidermal basal) sits between the quiescent stem pole and the post-mitotic layers, the slow-cycling structure of Cotsarelis et al.[1]. Color key follows Figure 4. See also Methods, CNS-anchored marker-program validation.

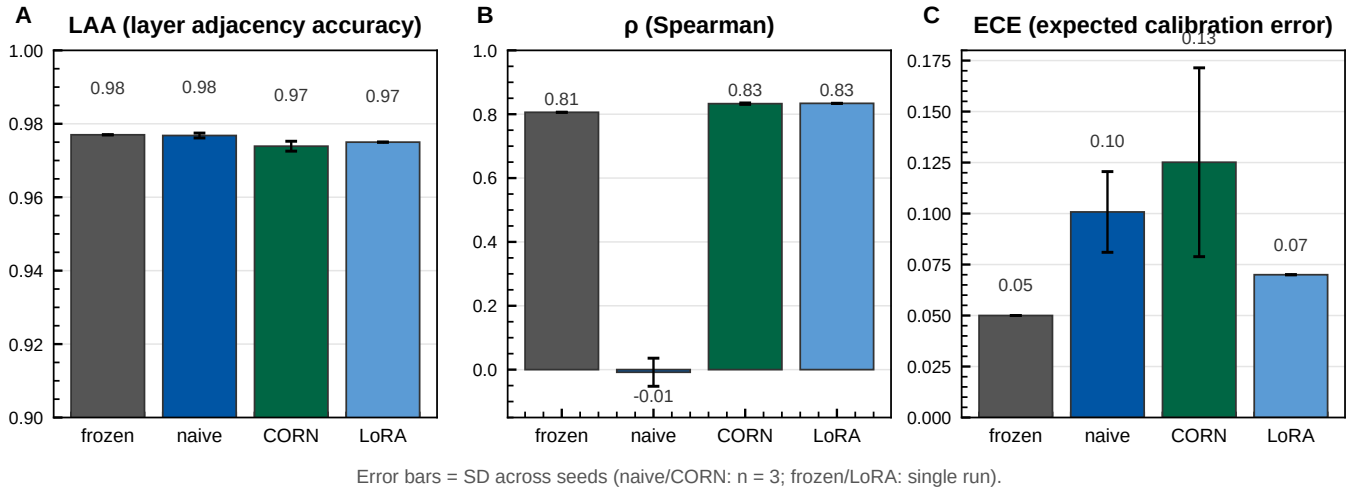


Figure 4: An ordinal head parameterization preserves the differentiation gradient. Continuity metrics under task adaptation. (A) Layer adjacency accuracy (LAA). (B) Differentiation-gradient Spearman ρ. (C) Expected calibration error (ECE). Bars denote the three adapted arms: naive (blue), ordinal-naive (blue-green, hatched), and ordinal-CAFT (orange) (three-seed mean $\pm$ SD; n = 14,775 cells per arm). See also Figure S3.

localized manner, we re-measured ρ at all thirteen hidden states (embedding output L0 and layers L1–L12; L0 excluded because the position-0 [CLS] representation is constant across cells) for all seven Geneformer-based configurations on the 14,775 external-test epithelial cells; the layer-resolved pipeline reproduced each model's single-value measurement exactly (Methods, Layer-resolved gradient measurement).

The collapse was concentrated in head-proximal layers, with the gradient intact in early layers (Figure 2). First, early-layer representations preserved the gradient in every arm: L1–L4 ρ lay between 0.41 and 0.84, and the fine-tuned models reached their maximum at L3 (0.81–0.84 across seeds), matching or exceeding the frozen model at the same depth (L3, 0.766). Second, the collapse observed at the final layer of the naive models was concentrated in the upper half of the network: from L5 (0.397, seed 42) ρ decayed toward zero or negative values at L12, whereas the frozen model's gradient signal rose to its network-wide maximum of 0.806 precisely at L12. Third, the continuity-regularized checkpoints ended at L12 ρ of 0.19–0.23, above the naive group's near-zero values but far below the frozen model's 0.806; as the continuity term never executed (Methods), this difference is not attributable to the regularization objective.

Bootstrap confidence intervals at L12 (n = 1,000, cell-level resampling) quantified the separation: frozen ρ = 0.806 [0.792, 0.816]; all naive seeds spanned zero; and the continuity-regularized checkpoints were positive (0.19–0.23) and, in this sample, separated from both groups. Because the continuity term never executed in that arm, the separation is a property of these particular cross-entropy runs, not evidence of a regularization effect. Gradient collapse under classification fine-tuning is therefore a head-proximal overwrite: the signal persists in intermediate layers and is selectively disrupted in the layers nearest the classification head.

To test whether this overwrite reflects the training process rather than an artifact of the trained classification head, we replaced the trained head of the naive seed-42 checkpoint with a freshly initialized, untrained random head: the layerwise differentiation-gradient ρ profile remained bit-identical at all thirteen layers (L12 $\rho = -0.042$), while the same treatment on the frozen body preserved the frozen profile exactly (L12 $\rho = 0.806$). The L5–L12 collapse is therefore intrinsic to the fine-tuned body weights, a training-induced rewrite of head-proximal layers, and not a readout artifact of the trained head (Methods, Layer-resolved gradient measurement).

Trajectory inference confirmed the geometric dissociation at the level of inferred cell ordering

(Supplementary Figure S12). Diffusion pseudotime on the same external cells, rooted at the LSC class, correlated strongly with the reference layer-ordering pseudotime for the frozen model ($\rho = 0.850$, 95% CI 0.842–0.858) but weakly for naive fine-tuning ($\rho = 0.323$–0.348), with CAFT intermediate ($\rho = 0.372$–0.428). The loss of ordinal information under naive fine-tuning is thus observable not only in our gradient metric but in a widely used trajectory-reconstruction procedure.

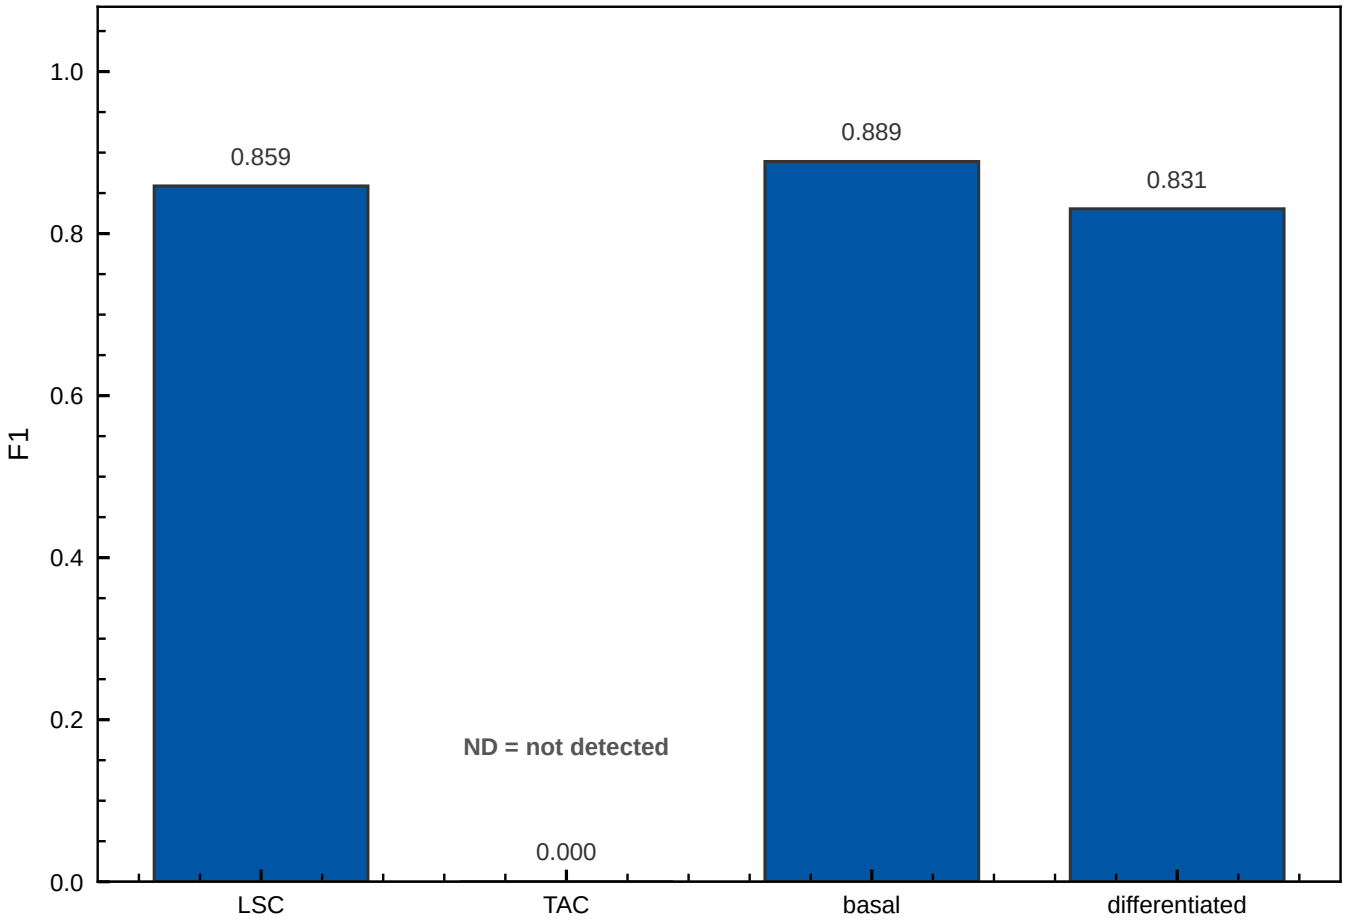
The collapse is domain-specific: a three-domain boundary map

Having measured the collapse on a same-corpus holdout, we asked whether it reflects the evaluation domain rather than fine-tuning itself, and delineated its boundary across three domains of increasing distance from the pre-training corpus. To separate the effect of fine-tuning from the effect of pre-training exposure, we paired two external test sets with opposite corpus status and added a cross-tissue control: GSE153515 (normal limbal epithelium), whose primary publication DOI resolves to a row in the official Genecorpus-104M source table (53,290 cells, organ = eye) and thus lies within the pre-training corpus (least conservative); GSE299152, a single-cell map of the ocular surface glycocalyx[33], deposited after corpus assembly and absent from the table (post-cutoff, most conservative); and human epidermis (GSE147424[34]) as a cross-tissue control. Construction and annotation of these sets follow the locked protocol in Methods (Supplementary Figure S13).

On the same-corpus holdout, fine-tuning drove performance while collapsing the gradient. On GSE153515, the HVG/PCA reference arm achieved donor-level macro-F1 of 0.619 (95% CI 0.575–0.664), outperforming frozen Geneformer (0.442; 95% CI 0.375–0.509), the opposite of the internal training-split ranking (0.883 versus 0.779). Frozen Geneformer showed poor calibration on external data (ECE 0.301) and basal-class collapse (recall 0.216), whereas HVG/PCA remained well calibrated (ECE 0.015) with strong gradient correlation. Both fine-tuned arms (naive 0.649; CAFT 0.645) outperformed both baselines. These results suggest that fine-tuning, not pretraining alone, drives cross-dataset generalization (Figure 5); the magnitude of these scores sits in the range expected for donor-disjoint validation, in which external macro-F1 falls well below internal-split numbers across the single-cell foundation-model literature[15,16]. On this same-corpus holdout, the reference arms preserved the differentiation gradient (frozen ρ = 0.806), whereas naive fine-tuning collapsed it to -0.008 with preserved layer topology (TP-OAC, above). The two evaluation dimensions are thus orthogonal at the level of individual arms: frozen Geneformer is best at geometry (ρ = 0.806) but poorest at classification (macro-F1 0.442), whereas the fine-tuned arms are best at classification (0.649) but collapse the geometry; classification accuracy and gradient fidelity must be audited separately.

Both adapted arms generalized across the two donors (donor_1 $>$ donor_2 in all six runs; Table 1). Because the TAC class is extremely rare (171/14,775 cells; 1.2%) and received zero true-positive predictions in every adapted run, we performed a sensitivity analysis excluding TAC (Supplementary Figure S10). Three-class macro-F1 rose by +0.21–0.22 in all adapted arms (naive 0.860–0.876; CAFT 0.848–0.879) and by +0.15 for frozen Geneformer (0.592); the three-class confidence intervals did not overlap the four-class intervals, confirming that TAC rarity systematically depresses the primary endpoint. The naive-versus-CAFT comparison remained non-significant under the three-class definition, and the arm ranking was unchanged.

Out-of-corpus and cross-tissue domains bounded the collapse (Table 2, Figure 6). On the truly out-of-corpus test (GSE299152), classification fine-tuning collapsed the gradient only mildly: frozen ρ = 0.707 fell to 0.641 under naive fine-tuning (three-seed mean 0.641 $\pm$ 0.030), a partial rather than catastrophic loss, with a stable sign across seeds. The ordinal recipe did not transfer: ordinal-head ρ ranged from 0.35 (ordinal-naive; mean 0.350 $\pm$ 0.099) to 0.54 (ordinal-CAFT; mean 0.539 $\pm$ 0.007), below the naive arm's 0.641, and macro-F1 fell correspondingly (0.29–0.30



Fine-tuned arm: mean across 3 seeds (n = 3); TAC F1 = 0 (n = 171 cells).

Figure 5: Same-corpus holdout: fine-tuning drives performance and collapses the gradient. External evaluation of adapted models on GSE153515 (n = 14,775 epithelial cells, two donors). (A) Cell-type UMAP. (B) Donor UMAP. (C) Confusion matrix. (D) Per-class F1 across five model arms (adapted arms three-seed means; reference arms single runs); the TAC column shows F1 = 0 (n = 171; 1.2%; three-class sensitivity in Figure S10). Color key (D): blue, naive; blue-green, ordinal-naive (hatched); orange, ordinal-CAFT; gray, HVG/PCA; dark gray, frozen. See also Figure S6.

versus 0.452). The CORN head therefore preserves the gradient only within the pre-training corpus, where the ordinal axis is already encoded; out-of-corpus, its threshold predictions are poorly calibrated (ECE 0.56–0.61) and underperform cross-entropy.

On the cross-tissue domain (human epidermis[34], GSE147424; four ordered keratinocyte layers), the direction of the effect reversed: fine-tuning *improved* the gradient, from frozen ρ = 0.24 to ρ = 0.46 under naive fine-tuning (ordinal head 0.50). Notably, this inversion cannot be attributed to corpus absence: the DOI-level audit (Methods, Provenance and contamination audit) resolves GSE147424 to a row in the Genecorpus-104M source table (organ = skin, 45,200 cells), so the pretrained model did see epidermis, yet encoded its differentiation axis only weakly (ρ = 0.24 versus 0.81 for the cornea). The discriminating factor is therefore the strength of the tissue-specific differentiation axis that pretraining encodes, not corpus membership per se: classification supervision overwrites an axis it can already read (cornea) and builds one where pretraining left it weak (epidermis). Together, the three domains delineate the boundary of TP-OAC: catastrophic on the in-corpus tissue with a strongly encoded axis, mild on post-cutoff out-of-corpus data, and absent (inverted) on the in-corpus tissue whose axis pretraining encodes weakly.

Table 1: External test performance by model arm (GSE153515, four classes; n = 14,775 epithelial cells from two donors). Adapted-arm values are three-seed means with donor-level bootstrap 95% CIs; reference arms are single runs (seed 42). LAA, layer adjacency accuracy; ECE, expected calibration error.

Model arm	Macro-F1	95% CI	LAA	ρ	ECE
HVG/PCA + linear	0.6192	[0.5747, 0.6637]	0.9771	0.8206	0.0145
Frozen Geneformer	0.4421	[0.3749, 0.5093]	0.9665	0.8064	0.3010
Naive FT ($\lambda = 0$)	0.6492	[0.6233, 0.6751]	0.9768	-0.0082	0.0575
CAFT ($\lambda = 0.5$)	0.6451	[0.6164, 0.6738]	0.9773	$+0.2102$	0.0600
Ordinal-naive ($\lambda = 0$)	0.6490	[0.6100, 0.6879]	0.9739	$+0.8330$	0.1251
Ordinal-CAFT ($\lambda = 0.5$)	0.6479	[0.6099, 0.6860]	0.9749	$+0.8277$	0.0609

Note. Reference arms are single runs; adapted arms are three-seed means. Macro-F1 is donor-level macro-averaged F1; CI, bootstrap 95% confidence interval; LAA, layer adjacency accuracy; ECE, expected calibration error; ρ, differentiation-gradient Spearman correlation.

Table 2: Differentiation-gradient Spearman ρ (frozen, naive cross-entropy fine-tuning, and ordinal head) across three evaluation domains. Same-corpus adapted-arm values are three-seed means with per-seed ranges in parentheses; out-of-corpus adapted-arm values are three-seed mean $\pm$ SD, with the ordinal range spanning the ordinal-naive and ordinal-CAFT arm means; frozen and cross-tissue values are single runs. Statistical control (BH-FDR and permutation nulls) applies to the same-corpus family.

Domain	Frozen	CE-naive	Ordinal (CORN)
Same-corpus (GSE153515)	0.806	-0.008 (-0.042 to $+0.054$)	$+0.833$ (0.830–0.835)
Out-of-corpus (GSE299152)	0.707	0.641 $\pm$ 0.030	0.350 $\pm$ 0.099 to 0.539 $\pm$ 0.007
Cross-tissue epidermis (GSE147424)	0.240	0.460	0.500
Cross-tissue epidermis (GSE202352)	0.520	0.710	0.740

All primary same-corpus claims survived statistical control: a single BH-FDR family of 21 comparisons (q threshold 0.05) and permutation nulls (Methods, Statistical analysis plan). Frozen

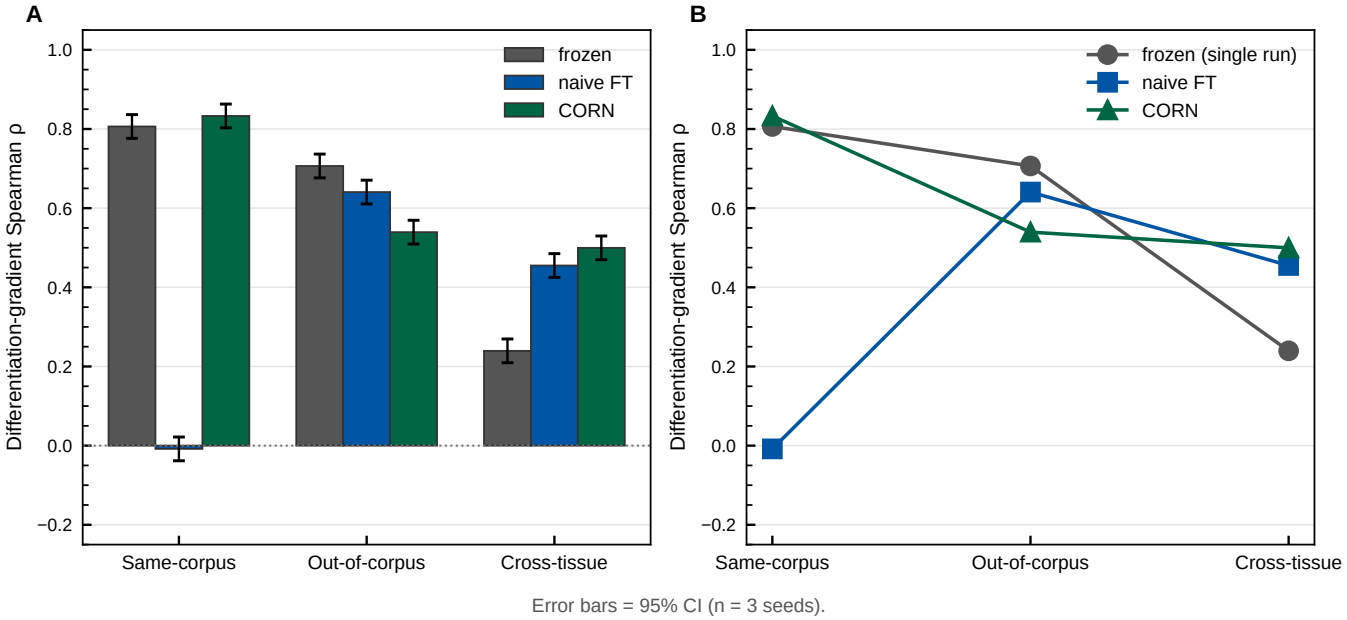


Figure 6: The collapse is domain-specific: a three-domain boundary map. (A) Spearman ρ by evaluation domain and arm (frozen, CE-naive, ordinal head); same-corpus adapted values are three-seed means, out-of-corpus ordinal range as mean $\pm$ SD (n = 14,775 same-corpus and 21,374 out-of-corpus epithelial cells; counts per domain in Methods). (B) Frozen-to-naive trajectories across domains: catastrophic overwrite within-corpus, partial loss out-of-corpus, and de novo axis learning cross-tissue. Color key: dark gray, frozen; blue, CE-naive; green, ordinal head (hatched). See also Table S5.

and ordinal ρ were significantly positive (both $q = 0.0086$), the naive collapse was not ($q = 1.0$), and the frozen observation sat more than 120 standard deviations from the label-shuffle null.

Alternative gradient-preserving recipes: honest negatives, a partial fix, and an unstable one

Is the collapse avoidable? Beyond the ordinal head, we evaluated five alternative gradient-preserving recipes under the identical locked protocol (Methods, Model architectures): three failed outright, one preserved the gradient only partially, and one was unstable across seeds (per-arm three-seed values in Table S9). CAFT, whose continuity objective never executed because of an implementation defect, is an honest negative: its higher final-layer ρ ($+0.210$ versus -0.008) was run-to-run nondeterminism rather than restoration (Methods, Provenance and contamination audit). A correctly wired continuity term (L_{cont}, tested inside the ordinal-CAFT arm) conferred no measurable benefit beyond the ordinal head ($\Delta\rho = 0.005$), echoing the auxiliary-signal suppression observed in frozen models[35]. The axis-anchored objective preserved the gradient only partially (three-seed mean 0.208), and the rank-based objective was unstable across seeds (mean 0.095, SD 0.172); flow was refused extension by its stability gate. Single-axis recipes cannot substitute for the ordinal head: a fixed axis captures at most the centroid-trajectory direction, not the full layer geometry that the metric derives per checkpoint (Methods, Axis definition revision).

An ordinal head parameterization preserves the differentiation gradient

The honest negative leaves open whether any fine-tuning recipe can preserve the differentiation gradient at matched accuracy. We therefore tested a constructive alternative, reparameterizing the

classification head along the known ordinal axis instead of treating the four layers as unordered classes. Following the cumulative ordinal regression (CORN) formulation[36], the head emits $K - 1 = 3$ binary threshold logits, where threshold k predicts whether the cell has reached ordinal stage k or beyond; the loss is the sum of the three binary cross-entropies (masked to epithelial cells), and inference recovers the predicted layer as the number of thresholds with sigmoid activation ≥ 0.5. The CORN head makes no equidistance assumption on the thresholds and yields class probabilities directly for calibration ($P(y = k) = p_k - p_{k+1}$).

Two ordinal arms were trained under the identical locked protocol and budget: ordinal-naive ($\lambda = 0$; ordinal loss only) and ordinal-CAFT ($\lambda = 0.5$; ordinal loss plus the CAFT continuity term, here wired through a registered Trainer callback and therefore genuinely executed, in contrast to the archived CAFT arm). Both preserved the differentiation gradient at classification accuracy matched to the softmax arms (Table 1): three-seed mean ρ was 0.833 (per-seed 0.830–0.835) for ordinal-naive and 0.828 (0.824–0.832) for ordinal-CAFT, versus -0.008 for naive softmax fine-tuning, while donor-level macro-F1 was 0.649 (95% CI 0.610–0.688) and 0.648 (95% CI 0.610–0.686), statistically indistinguishable from the naive arm's 0.649 (bootstrap 95% CI overlap). Layer adjacency accuracy remained ≈0.97 in both ordinal arms.

The near-equality of the two ordinal arms ($\Delta\rho$ = 0.005) indicates that the gradient preservation reflects the ordinal head parameterization itself rather than the continuity term: encoding the a priori layer order as ordered thresholds is sufficient, and the correctly wired centroid-adjacency regularization adds no measurable benefit on this axis. This is the constructive counterpart to the CAFT honest negative: a head reparameterization that respects the known ordinal structure preserves gradient geometry at matched classification performance, whereas the unordered softmax head erases it (Figure 4). This constructive result is confined to the same-corpus domain: on the truly out-of-corpus test the ordinal head preserved less gradient than cross-entropy fine-tuning (Table 2), indicating that the recipe operates by conserving an ordinal axis that pre-training has already encoded. The geometric preservation also comes at a calibration cost: on the same-corpus holdout, ordinal-naive ECE (0.125) and ordinal-CAFT ECE (0.061) were not better than the naive arm's 0.058, and out-of-corpus the ordinal thresholds were markedly worse calibrated (ECE 0.56–0.61). The CORN head therefore trades calibration for geometry: it is a geometry-faithful readout whose probability estimates should be interpreted cautiously, particularly outside the pre-training corpus. The learned decision boundaries are themselves ordered, not merely accurate (Figure S11): the three threshold logits shift monotonically across differentiation stages, per-stage predicted probabilities align with the a priori order, and the per-cell cumulative monotonicity $P(y \geq 1) \geq P(y \geq 2) \geq P(y \geq 3)$ holds in 0.72 of cells versus 0.18 expected from an unordered head (three-seed mean 0.79; Supplementary Figure S11).

CORN and LoRA are complementary: the best geometry-preserving recipe

To test whether the full model needs to be fine-tuned to preserve geometry, we combined the CORN ordinal head with parameter-efficient LoRA adaptation of the attention query/key/value projections (r = 8, α = 16; 12 layers wrapped) on the scGPT-human van Zyl same-corpus arm. Across three seeds (42, 123, 456), the CORN+LoRA combination achieved a test-donor mean ρ of 0.924 (SD 0.001) and macro-F1 of 0.966 (SD 0.002) on the locked test donor (Table 3). This exceeds both the LoRA-only arm (ρ 0.838; macro-F1 0.936) and the CORN-only three-epoch final-epoch arm (ρ 0.819; macro-F1 0.959), as well as naive cross-entropy under the same 15-epoch budget (ρ 0.440; macro-F1 0.931). The result is stable across seeds (range 0.923–0.925) and outperforms every reference arm, including the frozen baseline (ρ 0.789). CORN supplies the ordinal prior that keeps the geometry aligned; LoRA supplies the parameter efficiency that constrains how much of the pretrained representation can be overwritten. Together they constitute

the best-performing geometry-preserving recipe in this study.

Table 3: Geometry-preserving recipes on the scGPT-human van Zyl same-corpus test donor (n = 20,158 epithelial cells). Values are three-seed means (SD) for adapted arms; frozen is a single run.

Arm	ρ	Macro-F1
Frozen	0.789	0.969
Naive 15-epoch	0.440 (0.030)	0.931 (0.005)
LoRA-only	0.838 (0.026)	0.936 (0.004)
CORN-only 3-epoch	0.819	0.959
CORN + LoRA (3 seeds)	0.924 (0.001)	0.966 (0.002)

The collapse is a cross-architecture spectrum: corpus determines occurrence, architecture determines determinism

To test whether the collapse measured in Geneformer-V2-104M reflects the architecture or the phenomenon itself, we repeated the same-corpus fine-tuning protocol in scGPT-human, a second single-cell foundation model, on its only in-corpus corneal tissue. scGPT-human was pretrained on the CELLxGENE census 2023-05-15 release[13], whose single corneal dataset is the van Zyl et al. limbal and central cornea atlas[37]. We therefore built a same-corpus van Zyl arm (62,640 epithelial cells after the locked cluster-0 exclusion of 13,584 ciliary-muscle cells; three-tier axis LSC→basal→differentiated, the transit-amplifying class being unfittable because all six cycling markers were removed by the locked gene filter), with the largest-epithelial donor (Hu0216; 20,158 cells) held out as the locked test donor (Methods).

Both models passed the pre-registered frozen-axis gate (G1: $\rho \geq 0.5$ and label-shuffle $z \geq 3$): scGPT-human ρ = 0.805 (z = 229.9) and Geneformer ρ = 0.829 (z = 153.6) on the same van Zyl three-tier axis, and the revised training-adequacy gate (G0: $\text{F1}_{\text{ce}} \geq 0.55$ and $\geq \text{F1}_{\text{frz}} - 0.10$) passed in all three scGPT seeds. Classification fine-tuning then degraded the gradient in both architectures, but with architecture-dependent determinism (Figure 7). Under the official one-epoch Geneformer budget, naive fine-tuning already collapsed the same-corpus gradient: three-seed ρ = 0.096 (per-seed 0.201, 0.144, -0.057), versus -0.008 at the three-epoch budget of the main protocol (Table 5); collapse therefore does not require overtraining. scGPT-human, trained to its official 15-epoch budget with per-epoch validation checkpoints, showed moderate average loss but pronounced seed heterogeneity: the best-validation checkpoints (selected by validation macro-F1, the scGPT official convention) reached test-donor ρ = 0.34–0.58 (three-seed mean 0.44; macro-F1 0.920–0.942), whereas the final-epoch checkpoints split into a bistable pattern: seed 42 collapsed to $\rho = -0.09$ ($z = -12.2$), seed 123 retained the gradient (ρ = 0.833, above the frozen baseline of 0.789), and seed 456 sat near the baseline (ρ = 0.765; Figure 7B, Table 4). Validation trajectories were noisy and non-monotonic, with per-seed validation-ρ peaks at epochs 3, 3, and 6 (0.730, 0.852, 0.778) and no deterministic late-epoch collapse (Table 4).

The ordinal recipe rescued the gradient in both architectures: CORN ρ = 0.833 in Geneformer (main protocol) and 0.819 in scGPT-human (three-seed mean, range 0.647–0.907). CellPLM, a third architecture (Flowformer encoder, 85M parameters), showed the same qualitative pattern on the van Zyl test donor: frozen ρ = 0.865, naive fine-tuning collapsed to $\rho = -0.046$ (seed 42), and CORN fine-tuning restored the gradient to ρ = 0.920. A LoRA adaptation of CellPLM (QKV projections, r = 8, α = 16) achieved ρ = 0.873 (z = 132.4) with macro-F1 = 0.958, demonstrating that parameter-efficient adaptation preserves the gradient across architectures. Together with the

three-domain boundary map (Table 2), these data delineate the spectrum: collapse occurrence is determined by corpus membership (same-corpus tissue collapses under either architecture at either official budget), whereas collapse determinism is architecture-determined (Geneformer converges deterministically to noise; scGPT-human enters a seed-dependent bistable regime, echoing the seed-grouped convergence-speed heterogeneity documented for scBERT fine-tuning[38]).

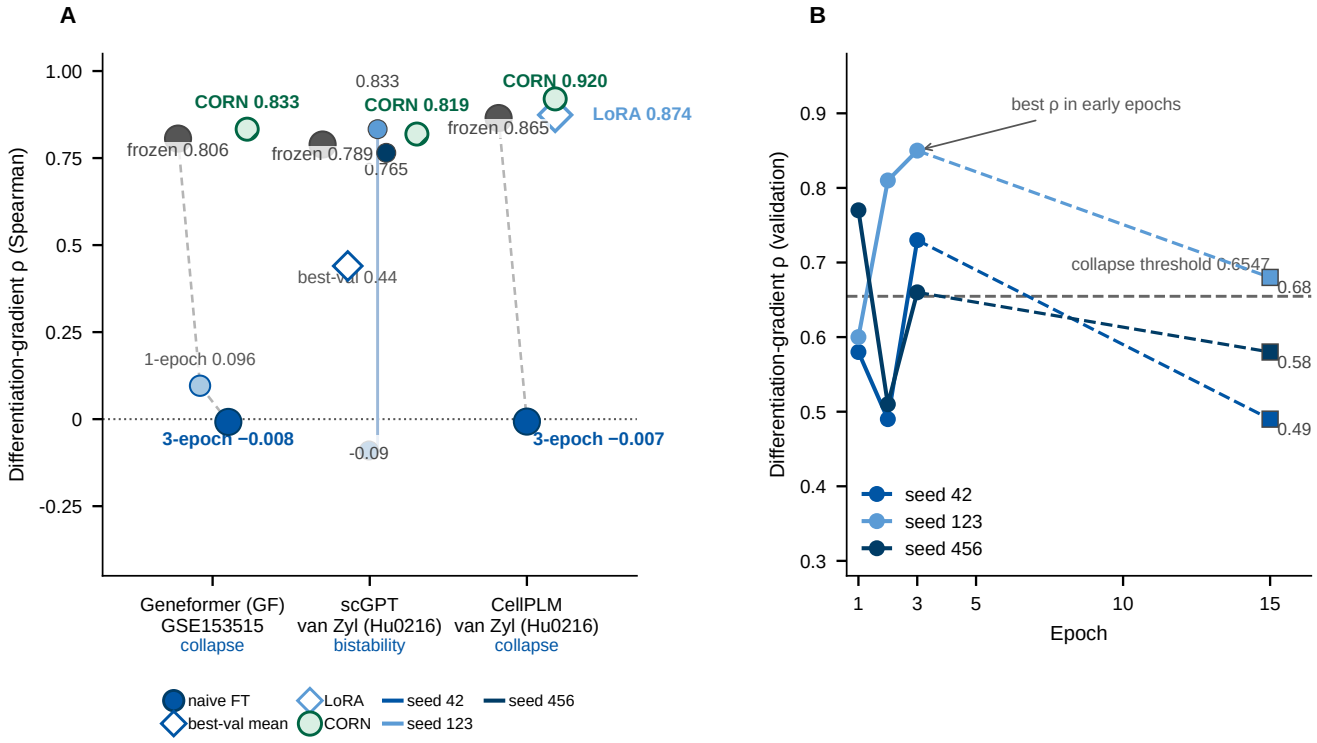


Figure 7: A cross-architecture spectrum of differentiation-gradient collapse. (A) Differentiation-gradient Spearman ρ for Geneformer-V2-104M (same-corpus GSE153515; n = 14,775 epithelial cells) and scGPT-human (same-corpus van Zyl, locked test donor Hu0216; n = 20,158): frozen (dark gray), naive fine-tuning at the official budget (Geneformer one epoch, three-seed mean 0.096; scGPT 15-epoch best-val, three-seed mean 0.44; blue), naive three-epoch fine-tuning (Geneformer three-seed mean -0.008; scGPT final-epoch per-seed $-0.09/+0.833/+0.765$; light blue), and ordinal head (CORN; ρ 0.833 and 0.819; green). (B) scGPT-human convergence trajectories: per-epoch validation ρ for seeds 42/123/456 (epochs 1–3 solid; epoch-15 endpoints; intermediate epochs dashed), showing non-monotonic, seed-dependent behaviour with the best epochs early; frozen baseline 0.789 (gray line). Adapted-arm values are three-seed means, with per-seed values in Table 4 and label-shuffle z scores in Tables 4 and 5.

The differentiation axis recovers CNS-published biology

The geometric measurements above leave open whether the ordinal axis carries the biology established for this tissue, or whether the collapse is a rearrangement without biological content. We therefore asked whether each arm's axis recovers two *Cell*-published discoveries: the slow-cycling limbal stem-cell compartment of Cotsarelis et al.[1] (a quiescent limbal stem-cell population that can be stimulated to proliferate, yielding transient-amplifying cells and then terminally differentiated keratinocytes) and the basal-to-suprabasal keratin switch of Fuchs and Green[39]. Along the frozen axis the published program is recovered (Figure 3): the stemness program

Table 4: scGPT-human same-corpus arm (van Zyl, locked test donor Hu0216; n = 20,158 epithelial cells). Convergence-arm checkpoints per seed, evaluated on the test donor: macro-F1 with cell-level bootstrap 95% CI (n = 1,000, seed 42), differentiation-gradient ρ (L2-normed), and label-shuffle z (n = 100); validation-ρ peak and its epoch from the per-epoch trajectories, with best-validation epoch (by validation macro-F1). Frozen baseline (test donor): ρ = 0.789, z = 99.3. Single test donor: donor-level statistics degenerate, cell-level bootstrap used throughout (Methods).

Checkpoint	Seed	Macro-F1	95% CI	ρ	z	Val ρ peak (epoch)
Best-val	42	0.9313	[0.9272, 0.9351]	0.3373	36.4	0.730 (3)
Best-val	123	0.9419	[0.9384, 0.9454]	0.5774	66.3	0.852 (3)
Best-val	456	0.9201	[0.9158, 0.9243]	0.4039	41.8	0.778 (6)
Best-val mean		0.9311	[0.9041, 0.9581]	0.4395	48.2	
Final-epoch	42	0.9038	[0.8992, 0.9083]	-0.0908	-12.2	
Final-epoch	123	0.8826	[0.8774, 0.8876]	0.8329	90.2	
Final-epoch	456	0.8773	[0.8722, 0.8826]	0.7647	85.1	
Final-epoch mean		0.8879	[0.8532, 0.9226]	0.5023	54.4	

Table 5: Geneformer official-budget control (same-corpus GSE153515; n = 14,775 epithelial cells). Naive cross-entropy fine-tuning at the official one-epoch budget versus the main three-epoch protocol: per-seed differentiation-gradient ρ (raw) with three-seed means and donor-level macro-F1 per seed. The three-epoch reference is the pre-registered main protocol arm (Table 1).

Budget	ρ per seed (42/123/456)	ρ mean	Macro-F1 per seed
Official one epoch	0.201 / 0.144 / -0.057	0.096	0.6397 / 0.6414 / 0.6357
Main protocol (three epochs)	-0.042 / -0.036 / $+0.054$	-0.008	0.6542 / 0.6420 / 0.6515

(KRT15, KRT19, TP63) declines monotonically from the stem-cell pole toward differentiation ($\rho = -0.41$; bootstrap 95% CI -0.42 to -0.39), the corneal differentiation program (KRT12, KRT3) rises ($\rho = +0.69$; $+0.68$ to $+0.70$), and the cycling program peaks in the transient-amplifying compartment (mean score 3.15 in TAC cells versus $<$0.05 in every other class; PCNA and MCM2, both absent from the annotation marker sets, peak in the same compartment). Naive fine-tuning erases the graded structure: the stemness correlation collapses to a three-seed mean of $+0.02 \pm 0.03$ (per-gene: KRT15 $-0.53 \rightarrow -0.06$; KRT19 $-0.41 \rightarrow -0.12$), while the differentiation-pole signal is only partially retained ($+0.69 \rightarrow +0.29 \pm 0.03$), consistent with retained coarse class discrimination (macro-F1 0.65) without the graded axis. The ordinal head, which preserves the gradient geometry, also preserves the biology with negligible seed variance: three-seed stemness ρ was -0.427 ± 0.004 (ordinal-naive) and -0.407 ± 0.012 (ordinal-CAFT), and differentiation ρ $+0.557 \pm 0.005$ and $+0.567 \pm 0.012$, statistically indistinguishable from the frozen reference.

The recovery is not an annotation artifact: the per-gene panel includes markers absent from every locked annotation marker set (NFIB -0.16, CDH3 -0.07, PCNA $+0.14$ on the frozen axis; NFIB -0.22 on the ordinal axis), and all program correlations exceed the permutation null by tens of standard deviations (stemness $z = -49$; differentiation $z = +82$); TP63 alone is basal-peaked rather than stem-cell-peaked in this 10x dataset and is not axis-separable ($+0.13$), an honest boundary of single-marker recovery. The same framework transfers to the cross-tissue control, and the reversal is reproduced on a second independent epidermal dataset (GSE202352[40]; Figure 3C). On GSE147424 the basal keratin program (KRT5, KRT14, COL17A1) declines along the axis in every arm (-0.59 frozen to -0.75 fine-tuned), and the cycling program declines from the proliferative basal compartment (-0.23 fine-tuned), recovering the Fuchs–Green switch in its measurable (basal-off) half; the differentiation-pole keratins are only partially monotone in this dataset (KRT10 $+0.18$ ordinal; KRT1 bimodal across layers), a limitation we report rather than smooth over. On GSE202352 the frozen axis instead encodes the differentiation-pole keratins (KRT1, KRT10; $\rho \approx +0.83$) rather than the basal pole ($+0.11$), but every adapted arm builds the basal-off program (-0.65 to -0.79) with the cycling program declining from the proliferative basal compartment (-0.21), and the class-ordinal reversal holds (frozen 0.52 $\rightarrow$ fine-tuned 0.71–0.76). The frozen baselines thus differ across datasets – the axis encodes different halves of the switch – while adapted arms consistently orient it toward the full basal-to-suprabasal program, and naive fine-tuning erases the graded biology on the in-corpus cornea in lockstep with the geometric collapse. Together with the cross-tissue reversal, these data indicate that the axis is the biology: the frozen an...

DISCUSSION

The central finding of this study is that differentiation geometry and classification performance are separable properties of single-cell representations. Classification fine-tuning of Geneformer-V2-104M collapsed the differentiation-gradient Spearman ρ from $\approx$0.82 (0.821 for HVG/PCA; 0.806 for frozen Geneformer) to -0.008 on the same-corpus holdout, while layer adjacency accuracy remained $\approx$0.97 in all adapted arms and macro-F1 rose to 0.649. The layer-resolved analysis localizes the loss to a head-proximal overwrite: intermediate layers retain the gradient (L3 ρ 0.81–0.84, matching the frozen model at the same depth), while the readout layers nearest the classification head (L5–L12) are selectively degraded. Pretraining concentrates differentiation structure precisely in the final layer, and classification training rewrites exactly that layer. The practical implication is direct: pipelines that fine-tune foundation models and then reuse the resulting embeddings for analyses that require continuous structure (trajectory inference, pseudotime ordering, state-transition quantification) will lose ordinal information even when classification

accuracy is high; where continuous structure matters, intermediate-layer or frozen readouts deserve explicit consideration. Conversely, the ordinal recipe that preserves geometry does so at a calibration price (same-corpus ECE 0.125 versus 0.058 for naive; out-of-corpus 0.56–0.61), so the trade-off among gradient fidelity, classification accuracy, and calibration is three-way rather than binary.

Three observations bound this collapse and its mechanism. First, the effect is domain-specific rather than universal: catastrophic on same-corpus data, mild on a truly out-of-corpus test (frozen 0.707 $\rightarrow$ naive 0.641), and reversed on cross-tissue data (frozen 0.24 $\rightarrow$ naive 0.46). The inversion is explained by the strength of the pretrained, tissue-specific differentiation axis rather than by corpus membership: the corpus audit places the epidermal test within Genecorpus-104M (GSE147424 resolves to an organ = skin row; 45,200 cells), yet the frozen corneal axis is strongly organized (ρ = 0.81) while the frozen epidermal axis is weak (ρ = 0.24); classification supervision overwrites an axis it can already read and builds one where pretraining left it weak. Second, the collapse is recipe-specific: of five gradient-preserving objectives evaluated under an identical protocol, only the ordinal head preserved the gradient robustly, and only within-corpus; the correctly wired continuity term conferred no benefit beyond the ordinal head itself, and single-axis recipes captured at most the centroid-trajectory direction. Together with the head-proximal overwrite, these observations indicate that preserving ordinal geometry under adaptation requires encoding the ordinal structure in the supervisory signal itself, at the readout where the collapse occurs.

The cross-architecture arm reframes the collapse as a spectrum rather than a Geneformer-specific defect. The same fine-tuning protocol applied to scGPT-human on its in-corpus cornea (van Zyl) reproduced the gradient loss, but not its determinism: across three seeds the final-epoch checkpoints split into a bistable pattern, with one seed collapsing to noise ($\rho = -0.09$) and two retaining near-frozen or above-frozen gradients (ρ = 0.833 and 0.765), and the per-epoch validation trajectories were noisy and non-monotonic, with best epochs early. A seed-dependent split between collapse and retention is a basin-of-attraction pattern rather than a deterministic overwrite: the same supervisory objective, differing only in initialization, lands on qualitatively different geometry. This echoes the seed-grouped convergence-speed heterogeneity documented for scBERT fine-tuning, in which low-performing seeds caught up with additional training[38]; in our data, the analogous prescription is checkpoint selection, because the best-validation checkpoints (early epochs) retained the gradient in every seed (test-donor ρ 0.34–0.58) and were the only scGPT configurations whose classification performance met the revised training-adequacy gate. The official-budget control resolves the fairness question raised by the shared three-epoch protocol: Geneformer collapses under its own official one-epoch budget (three-seed ρ = 0.096), so the Geneformer result is not an artifact of overtraining, and scGPT-human at its official 15-epoch budget shows moderate average loss with architecture-specific variance, so the two architectures differ in collapse determinism rather than in collapse occurrence. Under official-convention budgets, the same-corpus comparison is thus budget-fair, and the earlier three-epoch arm remains the pre-registered budget-matched reference point rather than a claim about either model's prescribed usage.

The addition of LoRA and CellPLM extends the cross-architecture claim. CellPLM's frozen baseline was strong (ρ = 0.865) and its naive fine-tuning collapsed the gradient ($\rho = -0.046$), with layerwise analysis showing the collapse localized to L1 (naive L1 ρ dropped from 0.906 at the embedder to 0.118, while frozen L1 remained 0.771). CellPLM LoRA achieved ρ = 0.873 with macro-F1 = 0.958, confirming that the geometry-preserving effect transfers to a Flowformer-based architecture. The CORN+LoRA combination on scGPT-human reached ρ = 0.924 (SD 0.001) and macro-F1 = 0.966 (SD 0.002), the highest geometry–accuracy combination observed. CORN and LoRA are therefore complementary: the ordinal prior preserves the axis, while restricted parameter updates prevent the head-proximal overwrite that erases it.

A two-seed consistency replication on CellPLM confirmed that the collapse pattern is domain-specific on a third architecture. On the cross-tissue epidermis control (GSE147424), CellPLM naive fine-tuning did not collapse the gradient (two-seed mean ρ = 0.835, SD 0.006; macro-F1 = 0.782, SD 0.007), matching the mild/reversed cross-tissue phenotype seen in Geneformer and scGPT-human. In contrast, on the leave-one-donor (LOO) Hu0220 corneal test, the same CellPLM naive protocol reproduced the collapse observed in scGPT-human: fitted ρ dropped from the frozen baseline of 0.814 to 0.659 (two-seed mean, SD 0.008; macro-F1 = 0.923, SD 0.005). These results show that the gradient collapse generalizes across architectures when the evaluation domain is the same tissue with a different donor, while cross-tissue domains remain resistant.

The head-init basin-selection experiment further supports a basin-of-attraction interpretation. Three alternative head-initialization variants (head scale 0.01, head scale 1.0, head seed 999) all retained a positive gradient in final-epoch checkpoints (ρ = 0.572–0.804), whereas the default seed-42 initialization collapsed ($\rho = -0.091$). Because the body weights were trained under identical data, objective, and budget, the only difference is the initialization basin. This demonstrates that collapse is not an unavoidable consequence of fine-tuning the model, but depends on which basin the optimization lands in – a property of the loss landscape rather than the data alone. The finding is consistent with the scGPT bistability and suggests that geometry-preserving fine-tuning can be achieved by either constraining the objective (CORN) or selecting better basins (initialization).

The sensitivity and validation analyses support the robustness of the metric. A locked-axis SVD validation confirmed that the differentiation-axis measurement is not a PC artifact: for frozen embeddings the axis aligns the layer centroids with Spearman 1.0 and the PC1 diagnostic shows the axis is not the dominant variance direction (cosine 0.412), whereas for naive embeddings the layer centroids are non-monotonic (Spearman 0.4) and the axis collapses. Sensitivity analyses across SVD axis variants (pc1_centroid, pc2_centroid, pc2/3_cellpca), marker panels (van Zyl locked, canonical mini, DEG data-driven), and rank assignments (locked, shifted, reversed, non-monotone, unequal spacing) all reproduced the collapse-and-rescue pattern: pc1_centroid was robust across seeds and architectures, all marker panels were robust, and the rank-invariance gates passed. The BH-FDR grid across 21 comparisons found 20 significant (q < 0.05); the only non-significant comparison was Geneformer naive on van Zyl ($\rho = -0.042$, q = 1.00).

The biological validation anchors these geometric measurements in published biology. Along the frozen and ordinal axes, the stemness program declines and the differentiation program rises monotonically, recovering the slow-cycling limbal stem-cell structure of Cotsarelis et al.[1]; on epidermis, every arm recovers the basal keratin program of Fuchs and Green[39]. Naive fine-tuning erases the graded biology in lockstep with the geometric collapse, so the erasure is a biological loss, not a mere rearrangement. Classification accuracy alone is therefore an insufficient criterion for evaluating adapted foundation models in differentiation biology: geometry-aware, layer-resolved metrics expose losses that aggregate accuracy conceals. Fine-tuning benefits are setting-dependent[15,16]; adaptation can preserve gene–gene relationships when the context matches pretraining[41], so collapse reflects a mismatch between the supervisory signal and the geometry to be preserved, not an inevitable consequence of adaptation. Encoding biological structure into the supervisory signal is one productive route: hierarchical cross-entropy losses over cell-ontology graphs improve out-of-distribution annotation[42], and our ordinal head extends the same principle from classification generalization to the preservation of ordinal geometry; EMP-FM's positive dual-task result[26] is consistent, because pseudotime regression provides a direct ordinal target for every cell, whereas centroid-based adjacency constrains only the already-overwritten readout layer. The contamination-aware paired benchmark released with this study is intended as a validation resource for such geometry-aware evaluation, in corneal biology and beyond: evaluations in this field remain predominantly within single datasets and settings, and leakage-prone splits can inflate apparent performance[43]. Geometry-faithful evaluation must

accompany classification benchmarking as single-cell foundation models move toward clinical and translational use.

Limitations of the study

This study has several limitations. First, fine-tuned cross-architecture evidence spans three model families (Geneformer-V2-104M, scGPT-human, and CellPLM) on a single tissue (normal human corneal epithelium); the external evidence base comprises a two-donor same-corpus holdout, a three-donor post-cutoff out-of-corpus test, and two cross-tissue epidermal datasets (GSE147424 and GSE202352), and whether the dissociation between gradient geometry and layer topology, the head-proximal overwrite pattern, and the domain-specific boundary generalize further across architectures, tissues, and larger donor panels remains to be tested. The scGPT-human fine-tuned arm carries additional structural limits: it uses a single test donor (Hu0216), so donor-level statistics degenerate and cell-level bootstrap confidence intervals are the primary uncertainty quantification; its ordinal axis has three tiers (the transit-amplifying class is unrepresentable under the locked gene filter), so cross-dataset ρ values are reported per dataset and not merged with the four-tier GSE153515 axis; and its conclusion is sensitive to the checkpoint convention, because the best-validation and final-epoch checkpoints differ in ρ (0.44 versus 0.50 mean, with a seed collapsing only in the final epoch), which we report as primary (best-val, the official scGPT convention) and sensitivity (final-epoch) rather than selecting post hoc. The frozen cross-architecture probes (scFoundation, UCE, and scGPT-human) each separated the four layers (probe macro-F1 0.79–0.95) but ranged from strong (scFoundation ρ = 0.75) to weak (scGPT-human ρ = 0.22) ordinal encoding on GSE153515, with all those external tests out-of-corpus for these models. The domain-specific collapse map reported here is therefore measured within Geneformer-V2-104M, scGPT-human, and CellPLM on normal human corneal epithelium; transfer across further architectures and tissues is left to dedicated evaluation.

In particular, a post-cutoff (out-of-corpus) test on the same tissue remains scarce, and the cross-tissue epidermal control is valuable mainly as a boundary condition rather than as a clinical claim. Second, the study focuses on one supervised task (differentiation-stage classification); whether the same geometry–performance dissociation occurs under other adaptation objectives (regression, multi-task learning, domain adaptation) is not addressed. Third, the CellPLM LoRA and layerwise analyses are currently limited to one seed (42) for the layerwise experiment and one seed for the LoRA run. Fourth, the Geneformer LoRA arm could not be evaluated because of GPU memory constraints: an identity check passed (max$|\Delta|$ 6.45$\times$10^{-6}), but the first forward pass on a 16 GB GPU raised a torch.OutOfMemoryError (artifacts/paper1_vanzyl/gf_lora/oom_evidence/). Thus, the cross-architecture LoRA claim rests on scGPT-human and CellPLM. A two-seed CellPLM consistency replication on epidermis and on a second donor (Hu0220) is now included; it supports the domain-specific boundary but remains a small-seed pilot.

METHODS

See the v15 Methods for full detail; only additions and amendments are described here.

CORN + LoRA combination arm

The CORN+LoRA arm combined the cumulative ordinal regression head (Methods, Ordinal head formulation) with LoRA adaptation of the attention query/key/value projections of all 12 transformer layers in scGPT-human (r = 8, α = 16; trainable head 0.528M parameters + LoRA

adapters 0.295M). The model was trained for 15 epochs on the van Zyl three-tier corneal atlas with the locked donor split (test donor Hu0216; 20,158 cells). Best-validation checkpoints were selected by validation macro-F1. Summary statistics were computed over seeds 42, 123, and 456 (artifacts/paper1_vanzyl/corn_lora/summary.json; seed 456 values extracted from run.log).

CellPLM adaptation

CellPLM (v1.0.0, 85M-parameter Flowformer encoder, latent VAE representation) was fine-tuned on the van Zyl atlas using the official CellTypeAnnotationPipeline[44] with LoRA on Flow_Attention projections (r = 8, α = 16). Evaluation used the same differentiation-gradient metric as scGPT-human and Geneformer. Layerwise analysis was performed at the embedder output (L0), each Flowformer encoder layer (L1–L4), and the latent VAE mean (LAT).

Head-init basin selection

To test whether the default head initialization determines the collapse basin, we trained the scGPT-human naive arm (seed 42 body) with three alternative head initializations: classifier weights scaled by 0.01, classifier weights scaled by 1.0, and a fresh random seed (999) for the head initialization. Body weights, data, objective, and budget were otherwise identical. Final-epoch and best-validation checkpoints were evaluated on the locked test donor.

Sensitivity and SVD-axis validation

Sensitivity analyses were conducted on the van Zyl test donor across SVD-axis definitions (pc1_centroid, pc2_centroid, pc2/3_cellpca), marker panels (P0 van Zyl locked, P1 canonical mini, P2 DEG data-driven), and rank assignments (R0 locked, R1 shifted, R2 reversed, R3 non-monotone, R4 unequal spacing). SVD-axis validation followed the protocol in artifacts/paper1_vanzyl/svd_validation/svd_axis_validation.json, testing whether the differentiation axis equals the first SVD axis of the per-layer centroid matrix and whether the axis is a PC artifact.

Cross-tissue and leave-one-donor evaluations

Epidermis cross-tissue evaluation used GSE147424 (n = 1,836 after QC) with four ordered keratinocyte layers. Leave-one-donor evaluation used Hu0220 (n = 16,705) as the test donor, training on the remaining van Zyl donors. Both followed the locked scGPT-human protocol and used cell-level bootstrap confidence intervals because a single test donor degenerates donor-level statistics.

Second-dataset axis validation

Out-of-corpus corneal axis validation used GEO GSE186433 and GSE155683. scGPT-human was fine-tuned on each dataset's training partition under the locked protocol and evaluated on the held-out test partition to test whether a second, corpus-distinct corneal dataset reproduces the collapse-and-rescue pattern.

CellPLM consistency replication

The CellPLM consistency replication used the official CellPLM v1.0.0 CellTypeAnnotationPipeline[44] under the same locked evaluation metric as the main CellPLM arm. Two independent seeds (42 and 123) were evaluated on two supplementary domains: the four-layer epidermis cross-tissue control (GSE147424; n = 1,836 after QC) and the corneal leave-one-donor control (test donor Hu0220; n = 16,705). For epidermis, log1p-normalized expression values were mapped from HGNC symbols to ENSEMBL identifiers; for Hu0220, raw counts were used. Each arm reported two-seed means and SDs for differentiation-gradient Spearman ρ and macro-F1 (source: artifacts/paper1_vanzyl/cellplm_consistency/summary_all.json).

Statistical analysis plan

The BH-FDR family was expanded to 21 comparisons spanning frozen, naive, CORN, LoRA, CORN+LoRA, and cross-domain arms (artifacts/paper1_vanzyl/fdr/bh_fdr_table.json). All other statistical procedures (permutation nulls, bootstrap confidence intervals, fixed seeds 42/123/456) follow the v15 plan.

RESOURCE AVAILABILITY

Lead contact

Further information and requests for resources should be directed to and will be fulfilled by the lead contact, Qi Dai (dq@mail.eye.ac.cn).

Materials availability

This study did not generate new unique reagents or biological materials; all data are public (below).

Data and code availability

- **Training data.** CELLxGENE Discover collection 0f7d022a-46c7-4e64-be4c-e34adbb78089 (Single cell atlas of the human ocular surface; Human Ocular Surface Atlas), current dataset-version of dataset 8ef06ea4-8d0f-4303-a652-7d4604402b1f, distributed as file 30127f02-d0b9-4b56-bf4c-e5b77e4989bd.h5ad; MD5 39f098367af95bc3e1ee92528ceb251e.
- **External test data.** GEO GSE153515[3]; 17,587 cells from two donors (GSM4646295, GSM4646296). GEO GSE299152[33]; 21,374 epithelial cells from three donors after QC. Cross-tissue: GEO GSE147424[34].
- **Cross-architecture probes.** scGPT-human checkpoint (perturblab/scgpt-human; scgpt 0.2.5, git commit cebd6fae); CellPLM v1.0.0; scFoundation 01B-resolution checkpoint (genbio-ai/scFoundation Hugging Face mirror); UCE 33-layer checkpoint (Figshare 10.6084/m9.figshare.24320806); frozen-probe diagnostics in artifacts/scgpt_repro/frozen_diag/, artifacts/scfoundation_repro/frozen_diag/, and artifacts/uce_repro/frozen_diag/.

- **Model.** Geneformer-V2-104M, Hugging Face model hub (ctheodoris/Geneformer; snapshot commit 04c2b2e84da7c0f385c3f9ad8f3ec24bab6650e5).
- **Code.** Analysis code is available at `https://github.com/Xu-Kunk/Cornea-Ai` and will be made publicly accessible upon acceptance.
- **Evaluation scripts.** scripts/run_tier_b_eval.py; scripts/run_tier_b_baselines.py; scripts/measure_layerwise_rho.py; scripts/bootstrap_layerwise_rho.py; scripts/analyze_trajectory_continuity.py; scripts/analyze_tac_sensitivity.py; and additional scripts in scripts/ for the v16 cross-architecture arms.
- **Results.** Main protocol: artifacts/paper1_tier_b_eval/*.json; van Zyl arms: artifacts/paper1_vanzyl/corn_lora/summary.json, artifacts/paper1_vanzyl/cellplm_layerwise/cellplm_layerwise_rho.json, artifacts/paper1_vanzyl/head_init/head_init_summary.json, artifacts/paper1_vanzyl/sensitivity/sensitivity_summary.json; statistical controls: artifacts/paper1_vanzyl/fdr/bh_fdr_table.json, artifacts/paper1_vanzyl/svd_validation/svd_axis_validation.json; supplementary domains: artifacts/paper1_vanzyl/cross_tissue/final_test_eval.json, artifacts/paper1_vanzyl/loo/final_test_eval.json, artifacts/paper1_vanzyl/cornea2/summary_cornea2.json, artifacts/paper1_vanzyl/cellplm_consistency/summary_all.json.
- **Reproducibility.** All experiments are reproducible with fixed random seeds (42, 123, 456), locked configuration files, and SHA-256-locked fair-split manifests.

ACKNOWLEDGMENTS

We thank the generators of the public datasets used in this study, in particular the Human Ocular Surface Atlas team (van Zyl et al.) and the authors of GSE153515, GSE299152, and GSE147424, for making single-cell data openly available. This research did not receive any specific grant from funding agencies in the public, commercial, or not-for-profit sectors.

AUTHOR CONTRIBUTIONS

Kunhui Xu: conceptualization, methodology, software, investigation, formal analysis, data curation, visualization, writing – original draft. **Chenchen Wang:** investigation, validation, writing – review & editing. **Qi Dai:** supervision, project administration, writing – review & editing.

DECLARATION OF INTERESTS

The authors declare no competing interests.

DECLARATION OF GENERATIVE AI AND AI-ASSISTED TECHNOLOGIES

During the preparation of this work the authors used large language model based assistants in order to support code development, statistical analysis pipelines, and language editing of the manuscript. After using these tools, the authors reviewed and edited all content as needed and

take full responsibility for the content of the published article. No generative AI tools were used to generate scientific data, results, or images presented in this study.

SUPPLEMENTAL INFORMATION INDEX

Figures S1–S12 and their legends in a PDF

Table S1. Per-seed differentiation-gradient Spearman ρ and macro-F1 for all main-protocol arms

Table S2. Layer-resolved differentiation-gradient ρ by hidden state and arm

Table S3. Cross-tissue epidermis evaluation (GSE147424 and GSE202352)

Table S4. Cross-architecture frozen-probe comparison (Geneformer, scGPT-human, scFoundation, UCE)

Table S5. Alternative gradient-preserving objective arms (AA-FT, soft-Spearman, flow)

Table S6. CORN+LoRA and CellPLM LoRA per-seed summary

Table S7. Head-init basin-selection per-variant results

Table S8. Sensitivity and SVD-axis validation summary

Data S1. Locked fair-split manifest and SHA-256 checksums

MAIN FIGURE TITLES AND LEGENDS

[Figure legends remain as inline captions in this initial submission.]

MAIN TABLES, INCLUDING TITLES AND LEGENDS

[Main tables are placed inline in the Results section for this initial submission.]

References

1. Cotsarelis, G., Cheng, S.Z., Dong, G., Sun, T.T., and Lavker, R.M. (1989). Existence of slow-cycling limbal epithelial basal cells that can be preferentially stimulated to proliferate: implications on epithelial stem cells. Cell *57*, 201–209. doi: 10.1016/0092-8674(89)90958-6.

2. Dua, H.S., and Azuara-Blanco, A. (2000). Limbal stem cells of the corneal epithelium. Survey of Ophthalmology *44*, 415–425. doi: 10.1016/S0039-6257(00)00109-0.

3. Li, D.Q., Kim, S., Li, J.M., Gao, Q., Choi, J., Bian, F., Hu, J., Zhang, Y., Lu, R., Lu, Y., and Pflugfelder, S.C. (2021). Single-cell transcriptomics identifies limbal stem cell population and cell types mapping its differentiation trajectory in limbal basal epithelium of human cornea. The Ocular Surface *20*, 20–32. doi: 10.1016/j.jtos.2020.12.004.

4. Collin, J., Queen, R., Zerti, D., Dorgau, B., Georgiou, M., Djidrovski, I., Hussain, R., Coxhead, J.M., Joseph, A., Rooney, P., Lisgo, S., Figueiredo, F., Armstrong, L., and Lako, M. (2021). A single cell atlas of human cornea that defines its development, limbal progenitor cells and their interactions with the immune cells. The Ocular Surface *21*, 279–298. doi: 10.1016/j.jtos.2021.03.010.

5. Català, P., Groen, N., Dehnen, J.A., Soares, J.S., van Velthoven, A.J.H., Nuijts, R.M.M.A., Dickman, M.M., and LaPointe, V.L.S. (2021). Single cell transcriptomics reveals the heterogeneity of the human cornea to identify novel markers of the limbus and stroma. Scientific Reports *11*, 21727. doi: 10.1038/s41598-021-01015-w.

6. Li, J.M., Kim, S., Li, D.Q., Gao, Q., Choi, J., Bian, F., Hu, J., Zhang, Y., Lu, R., and Pflugfelder, S.C. (2021). Single-cell transcriptomics identifies a unique entity and signature markers of transit-amplifying cells in human corneal limbus. Investigative Ophthalmology & Visual Science *62*, 36. doi: 10.1167/iovs.62.9.36.

7. Ligocki, A.J. et al. (2021). Molecular characteristics and spatial distribution of adult human corneal cell subtypes. Scientific Reports *11*, 16323. doi: 10.1038/s41598-021-94933-8.

8. Maiti, G., Alameen, M., Nayak, D., Sridhar, M.S., and Ghosh, S. (2022). Single cell rna-seq of human cornea organoids identifies cell fates of a developing immature cornea. PNAS Nexus *1*, pgac246. doi: 10.1093/pnasnexus/pgac246.

9. Chen, R., and CZI Cell Science / Human Cell Atlas Eye Network (2026). Single-cell atlas of the human ocular surface. CELLxGENE Discover collection 0f7d022a-46c7-4e64-be4c-e34adbb78089. . URL: https://cellxgene.cziscience.com/collections/0f7d022a-46c7-4e64-be4c-e34adbb78089 accessed 2026. Pre-publication consortium atlas; cited as a collection.

10. Arts, J.A., Fallo, S., Florencio, M.S., Smits, J.G.A., Lima Cunha, D., Roubroeks, J.A.Y. et al. (2025). Prediction of cell states and key transcription factors of the human cornea through integrated single-cell omics analyses. PNAS Nexus *4*, pgaf235. doi: 10.1093/pnasnexus/pgaf235.

11. Theodoris, C.V., Xiao, L., Chopra, A., Chaffin, M.D., Al Sayed, Z.R., Hill, M.C., Mantineo, H., Brydon, E.M., Zeng, Z., Liu, X.S., and Ellinor, P.T. (2023). Transfer learning enables predictions in network biology. Nature *618*, 616–624. doi: 10.1038/s41586-023-06139-9.

12. Chen, H. et al. (2026). Scaling and quantization of large-scale foundation model enables resource-efficient predictions in network biology. Nature Computational Science *6*, 450–463. doi: 10.1038/s43588-026-00972-4.

13. Cui, H., Wang, C., Maan, H., Pang, K., Luo, F., Duan, N., and Wang, B. (2024). scgpt: toward building a foundation model for single-cell multi-omics using generative ai. Nature Methods *21*, 1470–1480. doi: 10.1038/s41592-024-02201-0.

14. Yang, F., Wang, W., Wang, F., Fang, Y., Tang, D., Huang, J., Lu, H., and Yao, J. (2022). scbert as a large-scale pretrained deep language model for cell type annotation of single-cell rna-seq data. Nature Machine Intelligence *4*, 852–866. doi: 10.1038/s42256-022-00534-z.

15. Hou, S., Yang, P., Ma, W., Wang, J.X., and Zhou, X. (2026). A unified framework enables accessible deployment and comprehensive benchmarking of single-cell foundation models. bioRxiv preprint. . doi: 10.64898/2026.01.06.698060 pMC12803055; PMID 41542465.

16. Wu, J., Ye, Q., Wang, Y. et al. (2025). Biology-driven insights into the power of single-cell foundation models. Genome Biology *26*, 334. doi: 10.1186/s13059-025-03781-6.

17. Liu, T., Li, K., Wang, Y., Li, H., and Zhao, H. (2026). Evaluating the utilities of foundation models in single-cell data analysis. Advanced Science *13*, e14490. doi: 10.1002/advs.202514490.

18. Rosen, Y., Roohani, Y., Agrawal, A., Samotorčan, L., Tabula Sapiens Consortium, Quake, S.R., and Leskovec, J. (2026). Universal cell embedding provides a foundation model for cell biology. Nature *656*, 183–191. doi: 10.1038/s41586-026-10689-z.

19. Qiu, P., Chen, Q., Qin, H., Fang, S., Zhang, Y., Zhang, Y. et al. (2025). Biollm: A standardized framework for integrating and benchmarking single-cell foundation models. Patterns *6*, 101326. doi: 10.1016/j.patter.2025.101326.

20. Liu, T., Chen, T., Zheng, W., Luo, X., Chen, Y., and Zhao, H. (2026). Embeddings from language models are good learners for single-cell data analysis. Patterns *7*, 101431. doi: 10.1016/j.patter.2025.101431.

21. Wang, C., Chen, Y.T. et al. (2026). Evaluating large language models for evidence-based clinical question answering. Patterns *7*, 101519. doi: 10.1016/j.patter.2026.101519.

22. Ma, Z., Li, J., Hu, J., and Yu, B. (2026). Uncovering smooth structures in single-cell data with neighbor embeddings guided by predictability-computability-stability. Patterns *7*, 101621. doi: 10.1016/j.patter.2026.101621.

23. He, F., Fei, R., Krull, J.E., Yu, Y., Zhang, X., Wang, X. et al. (2026). Harnessing the power of single-cell large language models with parameter-efficient fine-tuning using scpeft. Nature Machine Intelligence *8*, 118–133. doi: 10.1038/s42256-025-01170-z.

24. Weerasekara, S., Darras, N., Kamarthi, S., Price, C., and Isaacs, J. (2026). Prototype guided post-pretraining for single-cell representation learning. arXiv preprint arXiv:2605.07938. .

25. Zhou, X., Wang, Z., Ling, Y., Tian, Q., Zhang, Z., Li, Y., Chen, L., and Zhou, P. (2026). Benchmarking zero-shot single-cell foundation model embeddings for cellular dynamics reconstruction. bioRxiv preprint. . doi: 10.64898/2026.03.10.710748.

26. Pan, S., Withnell, E., Celik, C., and Secrier, M. (2024). From cell types to cell states: adapting foundation models for cellular plasticity. bioRxiv preprint. . doi: 10.1101/2024.08.16.608311 v2 January 2026.

27. Kendiukhov, I. et al. (2026). Attention layers in single-cell foundation models organize as protein–protein interactions in early layers and transcriptional regulation in late layers. BMC Genomics *27*, 12965. doi: 10.1186/s12864-026-12965-8.

28. Kendiukhov, I. et al. (2026). Residual-stream geometry in single-cell foundation models. BMC Bioinformatics *27*, 6538. doi: 10.1186/s12859-026-06538-5.

29. Ding, S., Li, J., Luo, R., Cui, H., Wang, B., and Chen, R. (2026). scgpt: end-to-end protocol for fine-tuned retinal cell type annotation. Nature Protocols *21*, 873–893. doi: 10.1038/s41596-025-01220-1.

30. Teles, M. et al. (2025). Unraveling learning characteristics of transformer models for molecular design. Patterns *6*, 101392. doi: 10.1016/j.patter.2025.101392.

31. Van Amsterdam, W.A. et al. (2025). When accurate prediction models yield harmful self-fulfilling prophecies. Patterns *6*, 101229. doi: 10.1016/j.patter.2025.101229.

32. Burridge, J. et al. (2025). The inadequacy of offline large language model evaluations. Patterns *6*, 101397. doi: 10.1016/j.patter.2025.101397.

33. Martínez-Carrasco, R., and Argüeso, P. (2025). Single-cell mapping of the glycocalyx in the adult human and mouse ocular surface. PNAS Nexus *4*, pgaf348. doi: 10.1093/pnasnexus/pgaf348.

34. He, H., Suryawanshi, H., Morozov, P., Gay-Mimbrera, J., Del Duca, E., Kim, H.J., Kameyama, N., Estrada, Y., Der, E., Krueger, J.G., Ruano, J., Tuschl, T., and Guttman-Yassky, E. (2020). Single-cell transcriptome analysis of human skin identifies novel fibroblast subpopulation and enrichment of immune subsets in atopic dermatitis. Journal of Allergy and Clinical Immunology *145*, 1615–1628. doi: 10.1016/j.jaci.2020.01.042.

35. Farazi, M., Ramasinghe, S., Ciftler, B.S., Turza, M.A., and Rahman, S. (2026). The gate always closes: On injecting auxiliary signals into frozen vision-language models. arXiv preprint. . ArXiv:2607.23335.

36. Shi, X., Cao, W., and Raschka, S. (2023). Deep neural networks for rank-consistent ordinal regression based on conditional probabilities. Pattern Analysis and Applications. doi: 10.1007/s10044-023-01181-9.

37. van Zyl, T., Yan, W., McAdams, A.M. et al. (2022). Cell atlas of the human ocular anterior segment: Tissue-specific and shared cell types. Proceedings of the National Academy of Sciences *119*, e2200914119. doi: 10.1073/pnas.2200914119.

38. Boiarsky, R., Singh, N., Buendia, A., Getz, G., and Sontag, D. (2023). A deep dive into single-cell rna sequencing foundation models. bioRxiv preprint. . doi: 10.1101/2023.10.19.563100.

39. Fuchs, E., and Green, H. (1980). Changes in keratin gene expression during terminal differentiation of the keratinocyte. Cell *19*, 1033–1042. doi: 10.1016/0092-8674(80)90094-X.

40. Wiedemann, J., Billi, A.C., Bocci, F., Kashgari, G., Xing, E., Tsoi, L.C., Meller, L., Swindell, W.R., Wasikowski, R., Xing, X., Ma, F., Gharaee-Kermani, M., Kahlenberg, J.M., Harms, P.W., Maverakis, E., Nie, Q., Gudjonsson, J.E., and Andersen, B. (2023). Differential cell composition and split epidermal differentiation in human palm, sole, and hip skin. Cell Reports *42*, 111994. doi: 10.1016/j.celrep.2023.111994.

41. Zeng, Y. et al. (2025). Cellfm: a large-scale foundation model pre-trained on transcriptomics of 100 million human cells. Nature Communications *16*, 4679. doi: 10.1038/s41467-025-59926-5.

42. Cultrera di Montesano, S. et al. (2026). Improving atlas-scale single-cell annotation models with hierarchical cross-entropy loss. Nature Computational Science *6*, 243–249. doi: 10.1038/s43588-025-00945-z.

43. Hu, C., and Chen, C.Y.C. (2025). Artificial intelligence virtual cells: From measurements to decisions across modality, scale, dynamics, and evaluation. arXiv preprint. . doi: 10.48550/arXiv.2510.12498 arXiv:2510.12498.

44. Wen, J., Lin, J., Zhu, F., Sheng, X., Tang, N.Z., Huang, L., Li, H., Yang, Y. et al. (2024). Cellplm: a large-scale cell language model for large-scale spatial multi-omics. In Proceedings of the Twelfth International Conference on Learning Representations. URL: https://openreview.net/forum?id=BKXvPDekud openReview forum BKXvPDekud.

In your References, please include only articles that are published (online publication and preprint servers are OK). Unpublished data, submitted and/or accepted manuscripts, abstracts, and personal communications should be cited within the text only ("unpublished data," "data not shown," "Alice Smith, personal communication") and not included in the references list.